



 This image cannot currently be displayed.

Chapter 2: Intro to Relational Model

Database System Concepts, 6th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use



Example of a Relation

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 12121 | Wu | Finance | 90000 |
| 15151 | Mozart | Music | 40000 |
| 22222 | Einstein | Physics | 95000 |
| 32343 | El Said | History | 60000 |
| 33456 | Gold | Physics | 87000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 58583 | Califieri | History | 62000 |
| 76543 | Singh | Finance | 80000 |
| 76766 | Crick | Biology | 72000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 98345 | Kim | Elec. Eng. | 80000 |

attributes
(or columns)

tuples
(or rows)



Attribute Types

- The set of allowed values for each attribute is called the **domain** of the attribute
- Attribute values are (normally) required to be **atomic**; that is, indivisible
- The special value ***null*** is a member of every domain. Indicated that the value is “unknown”
- The null value causes complications in the definition of many operations



Relation Schema and Instance

- A_1, A_2, \dots, A_n are *attributes*
- $R = (A_1, A_2, \dots, A_n)$ is a *relation schema*

Example:

instructor = (*ID*, *name*, *dept_name*, *salary*)

- Formally, given sets D_1, D_2, \dots, D_n a **relation** r is a subset of
 $D_1 \times D_2 \times \dots \times D_n$

Thus, a relation is a set of n -tuples (a_1, a_2, \dots, a_n) where each $a_i \in D_i$

- The current values (**relation instance**) of a relation are specified by a table
- An element t of r is a *tuple*, represented by a *row* in a table



Relations are Unordered

- Order of tuples is irrelevant (tuples may be stored in an arbitrary order)
- Example: *instructor* relation with unordered tuples

| <i>ID</i> | <i>name</i> | <i>dept_name</i> | <i>salary</i> |
|-----------|-------------|------------------|---------------|
| 22222 | Einstein | Physics | 95000 |
| 12121 | Wu | Finance | 90000 |
| 32343 | El Said | History | 60000 |
| 45565 | Katz | Comp. Sci. | 75000 |
| 98345 | Kim | Elec. Eng. | 80000 |
| 76766 | Crick | Biology | 72000 |
| 10101 | Srinivasan | Comp. Sci. | 65000 |
| 58583 | Califieri | History | 62000 |
| 83821 | Brandt | Comp. Sci. | 92000 |
| 15151 | Mozart | Music | 40000 |
| 33456 | Gold | Physics | 87000 |
| 76543 | Singh | Finance | 80000 |

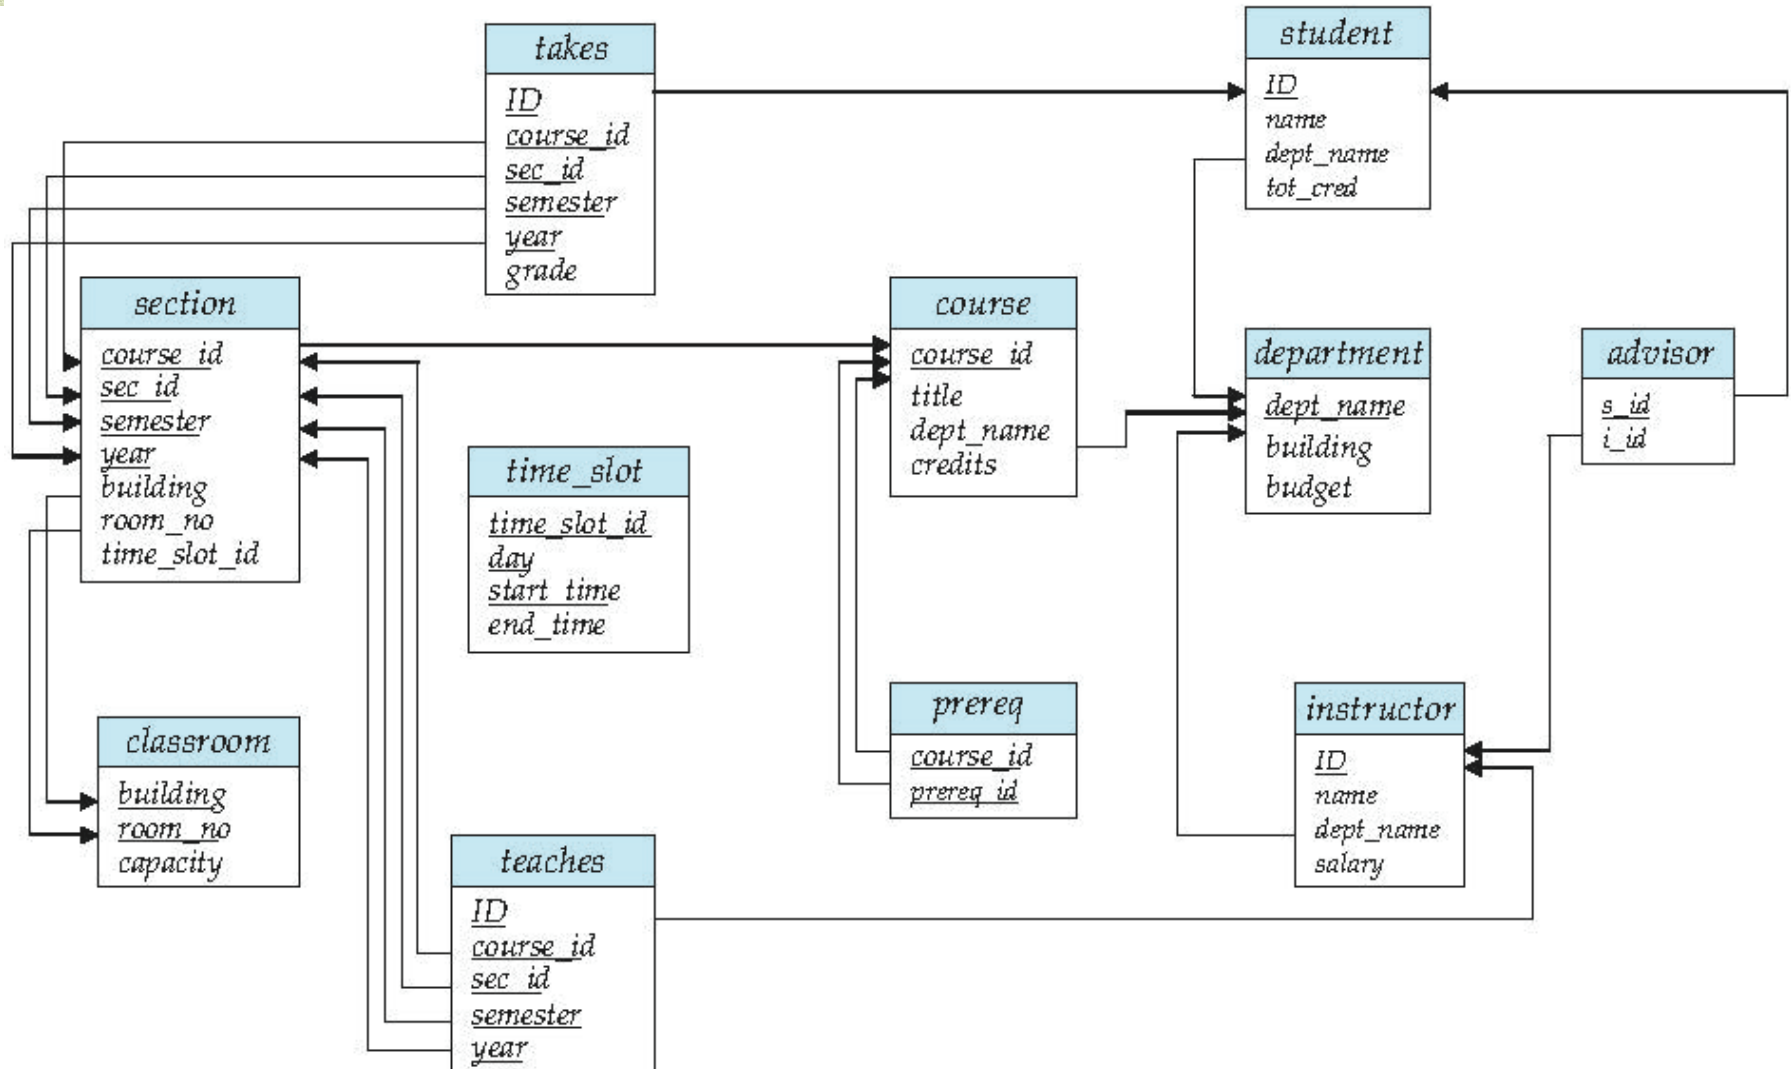


Keys

- Let $K \subseteq R$
- K is a **superkey** of R if values for K are sufficient to identify a unique tuple of each possible relation $r(R)$
 - Example: $\{ID\}$ and $\{ID, name\}$ are both superkeys of *instructor*.
- Superkey K is a **candidate key** if K is minimal
Example: $\{ID\}$ is a candidate key for *Instructor*
- One of the candidate keys is selected to be the **primary key**.
 - which one?
- **Foreign key** constraint: Value in one relation must appear in another
 - **Referencing** relation
 - **Referenced** relation
 - Example – *dept_name* in *instructor* is a foreign key from *instructor* referencing *department*



Schema Diagram for University Database





Relational Query Languages

- Procedural vs .non-procedural, or declarative
- “Pure” languages:
 - Relational algebra
 - Tuple relational calculus
 - Domain relational calculus
- The above 3 pure languages are equivalent in computing power
- We will concentrate in this chapter on relational algebra
 - Not turning-machine equivalent
 - consists of 6 basic operations



Select Operation – selection of rows (tuples)

- Relation r

| A | B | C | D |
|----------|----------|----|----|
| α | α | 1 | 7 |
| α | β | 5 | 7 |
| β | β | 12 | 3 |
| β | β | 23 | 10 |

- $?_{A=B \wedge D > 5}(r)$

| A | B | C | D |
|----------|----------|----|----|
| α | α | 1 | 7 |
| β | β | 23 | 10 |



Project Operation – selection of columns (Attributes)

- Relation r :

| A | B | C |
|----------|----|---|
| α | 10 | 1 |
| α | 20 | 1 |
| β | 30 | 1 |
| β | 40 | 2 |

- $\pi_{A,C}(r)$

| A | C |
|----------|---|
| α | 1 |
| α | 1 |
| β | 1 |
| β | 2 |

 $=$

| A | C |
|----------|---|
| α | 1 |
| β | 1 |
| β | 2 |



Union of two relations

- Relations r, s :

| A | B |
|----------|-----|
| α | 1 |
| α | 2 |
| β | 1 |

r

| A | B |
|----------|-----|
| α | 2 |
| β | 3 |

s

- $r \cup s$:

| A | B |
|----------|-----|
| α | 1 |
| α | 2 |
| β | 1 |
| β | 3 |



Set difference of two relations

- Relations r , s :

| A | B |
|----------|-----|
| α | 1 |
| α | 2 |
| β | 1 |

r

| A | B |
|----------|-----|
| α | 2 |
| β | 3 |

s

- $r - s$:

| A | B |
|----------|-----|
| α | 1 |
| β | 1 |



Set intersection of two relations

■ Relation r, s :

| A | B |
|----------|-----|
| α | 1 |
| α | 2 |
| β | 1 |

r

| A | B |
|----------|-----|
| α | 2 |
| β | 3 |

s

■ $r \bowtie s$

| A | B |
|----------|-----|
| α | 2 |

Note: $r \bowtie s = r - (r - s)$



joining two relations -- Cartesian-product

■ Relations r , s :

| A | B |
|----------|-----|
| α | 1 |
| β | 2 |

r

| C | D | E |
|----------|-----|-----|
| α | 10 | a |
| β | 10 | a |
| β | 20 | b |
| γ | 10 | b |

s

■ $r \times s$:

| A | B | C | D | E |
|----------|-----|----------|-----|-----|
| α | 1 | α | 10 | a |
| α | 1 | β | 10 | a |
| α | 1 | β | 20 | b |
| α | 1 | γ | 10 | b |
| β | 2 | α | 10 | a |
| β | 2 | β | 10 | a |
| β | 2 | β | 20 | b |
| β | 2 | γ | 10 | b |



Cartesian-product – naming issue

■ Relations r , s :

| A | B |
|----------|-----|
| α | 1 |
| β | 2 |

r

| B | D | E |
|----------|-----|-----|
| α | 10 | a |
| β | 10 | a |
| β | 20 | b |
| γ | 10 | b |

s

■ $r \times s$:

| A | $r.B$ | $s.B$ | D | E |
|----------|-------|----------|-----|-----|
| α | 1 | α | 10 | a |
| α | 1 | β | 10 | a |
| α | 1 | β | 20 | b |
| α | 1 | γ | 10 | b |
| β | 2 | α | 10 | a |
| β | 2 | β | 10 | a |
| β | 2 | β | 20 | b |
| β | 2 | γ | 10 | b |



Renaming a Table

- Allows us to refer to a relation, (say E) by more than one name.

$$\bowtie_X(E)$$

returns the expression E under the name X

- Relations r

| A | B |
|----------|-----|
| α | 1 |
| β | 2 |

r

- $r \bowtie \bowtie_S(r)$

| $r.A$ | $r.B$ | $s.A$ | $s.B$ |
|----------|-------|----------|-------|
| α | 1 | α | 1 |
| α | 1 | β | 2 |
| β | 2 | α | 1 |
| β | 2 | β | 2 |



Composition of Operations

- Can build expressions using multiple operations

- Example: $\pi_{A=C}(r \times s)$

- $r \times s$

| A | B | C | D | E |
|----------|---|----------|----|---|
| α | 1 | α | 10 | a |
| α | 1 | β | 10 | a |
| α | 1 | β | 20 | b |
| α | 1 | γ | 10 | b |
| β | 2 | α | 10 | a |
| β | 2 | β | 10 | a |
| β | 2 | β | 20 | b |
| β | 2 | γ | 10 | b |

- $\pi_{A=C}(r \times s)$

| A | B | C | D | E |
|----------|---|----------|----|---|
| α | 1 | α | 10 | a |
| β | 2 | β | 10 | a |
| β | 2 | β | 20 | b |



Joining two relations – Natural Join

- Let r and s be relations on schemas R and S respectively. Then, the “natural join” of relations R and S is a relation on schema $R \bowtie S$ obtained as follows:
 - Consider each pair of tuples t_r from r and t_s from s .
 - If t_r and t_s have the same value on each of the attributes in $R \bowtie S$, add a tuple t to the result, where
 - ▶ t has the same value as t_r on r
 - ▶ t has the same value as t_s on s



Natural Join Example

- Relations r, s :

| A | B | C | D |
|----------|-----|----------|-----|
| α | 1 | α | a |
| β | 2 | γ | a |
| γ | 4 | β | b |
| α | 1 | γ | a |
| δ | 2 | β | b |

r

| B | D | E |
|-----|-----|------------|
| 1 | a | α |
| 3 | a | β |
| 1 | a | γ |
| 2 | b | δ |
| 3 | b | ϵ |

s

- Natural Join

- $r \bowtie s$

| A | B | C | D | E |
|----------|-----|----------|-----|----------|
| α | 1 | α | a | α |
| α | 1 | α | a | γ |
| α | 1 | γ | a | α |
| α | 1 | γ | a | γ |
| δ | 2 | β | b | δ |

$$\circlearrowleft A, r.B, C, r.D, E \left(?_{r.B = s.B \wedge r.D = s.D} (r \times s) \right)$$



Notes about Relational Languages

- Each Query input is a table (or set of tables)
- Each query output is a table.
- All data in the output table appears in one of the input tables
- Relational Algebra is not Turing complete
- Can we compute:
 - SUM
 - AVG
 - MAX
 - MIN



Summary of Relational Algebra Operators

| Symbol (Name) | Example of Use |
|---------------------------------|---|
| σ (Selection) | $\sigma \text{ salary} \geq 85000$ (<i>instructor</i>) |
| | Return rows of the input relation that satisfy the predicate. |
| Π (Projection) | $\Pi ID, salary$ (<i>instructor</i>) |
| | Output specified attributes from all rows of the input relation. Remove duplicate tuples from the output. |
| \times (Cartesian Product) | <i>instructor</i> \times <i>department</i> |
| | Output pairs of rows from the two input relations that have the same value on all attributes that have the same name. |
| \cup (Union) | $\Pi name$ (<i>instructor</i>) \cup $\Pi name$ (<i>student</i>) |
| | Output the union of tuples from the <i>two</i> input relations. |
| $-$ (Set Difference) | $\Pi name$ (<i>instructor</i>) $--$ $\Pi name$ (<i>student</i>) |
| | Output the set difference of tuples from the two input relations. |
| \bowtie (Natural Join) | <i>instructor</i> \bowtie <i>department</i> |
| | Output pairs of rows from the two input relations that have the same value on all attributes that have the same name. |



 This image cannot currently be displayed.

End of Chapter 2

Database System Concepts, 6th Ed.

©Silberschatz, Korth and Sudarshan

See www.db-book.com for conditions on re-use