

SQL Programming

- SQL in Application Programs
- JDBC: Java Database Connectivity
- CLI: Call-Level Interface
- Embedded SQL

SQL in Applications Programs

- We have seen how SQL is used at the generic query interface --- an environment where we sit at a terminal and ask queries of a database.
- Reality is almost always different: conventional programs interacting with SQL.
- Want to consider:
 - How do we enable a database to interact with an “ordinary” program written in a language such as C or Java?
 - How do we deal with the differences in data types supported by SQL and conventional languages?
 - ▶ In particular, relations, which are the result of queries, are not directly supported by conventional languages.

Three-Tier Architecture

- A common environment for using a database has three tiers of processors:
 1. *Web servers* --- Connect users to the database, usually over the Internet, or possibly a local connection.
 2. *Application servers* --- Execute the “business logic” – whatever it is the system is supposed to do.
 3. *Database servers* --- Run the DBMS and execute queries and modifications at the request of the application servers.

Example: Amazon


- Database holds the information about products, customers, etc.
- Business logic includes things like “what do I do after someone clicks ‘checkout’?”
 - **Answer:** Show the “How will you pay for this?” screen.

For this section, we will deal with the interaction between the application and the DBMS

Environments, Connections, Queries

- A SQP **environment** is the framework under which data exists and SQL operations are executed.
- Think of a SQL environment as a DBMS running at some installation.
 - ◆ So tables, triggers, views, etc are defined within a SQL environment
- Database servers maintain some number of **connections**, so app servers can ask queries or perform modifications.
- The app server issues **statements**: queries and modifications, usually.

Diagrammatically



The SQL/Host Language Interface: Options

1. Code in a specialized language is stored in the database itself (e.g., PSM, PL/SQL).
 - ◆ Not covered (see text for info)
2. Connection tools are used to allow a conventional language to access a database (e.g. JDBC, CLI, PHP/DB).
3. SQL statements are embedded in a *host language* (e.g., C).

The Impedance Mismatch Problem

Basic problem: **impedance mismatch** – the data model of SQL differs significantly from the models of other languages.

- SQL uses the relational data model
- C, Java, etc., use a data model with ints, reals, pointers, records, etc.
- Consequently, passing data between SQL and a host language is not straightforward.

Host/SQL Interfaces Via Libraries

- The first approach to connecting databases to conventional languages is to use **library calls**.
 - Java + JDBC
 - C + CLI

SQL Programming: JDBC

- *Java Database Connectivity* (JDBC) is a library similar to SQL/CLI, but with Java as the host language.
- JDBC supports
 - Establishing a connection
 - Creating JDBC statements
 - Executing SQL statements
 - Getting a ResultSet
 - Closing connection

Making a Connection

Three initial steps:

1. Include

```
import java.sql.*;
```

to make the JAVA classes available.

2. Load a (vendor specific) “driver” for the database system being used.

```
Class.forName(“com.microsoft.sqlserver.jdbc.SQLServerDriver”);
```

dynamically loads a driver class for SQL Server db.

3. Establish a connection to the database.

```
Connection con = DriverManager.getConnection(“jdbc:mysql://  
localhost/Food? User=UserName&password=Password”);
```

establishes connection to database (Food) by obtaining a *Connection* object.

Making a Connection

```
import java.sql.*;  
Class.forName("com.microsoft.sqlserver.jdbc.SQLServerDriver");  
Connection myCon =  
    DriverManager.getConnection(<URL, name, passwd, etc>);
```

The JDBC classes

The driver For SQL Server. (Others exist)

Get an object of class Connection, which we've called myCon

The diagram illustrates the steps to make a JDBC connection. It features four code snippets: 'import java.sql.*;', 'Class.forName("com.microsoft.sqlserver.jdbc.SQLServerDriver");', 'Connection myCon =', and 'DriverManager.getConnection(<URL, name, passwd, etc>;'. Annotations with arrows point to these snippets: 'The JDBC classes' points to the import statement; 'The driver For SQL Server. (Others exist)' points to the Class.forName statement; and 'Get an object of class Connection, which we've called myCon' points to the DriverManager.getConnection statement.

Statements

- JDBC provides two classes:
 1. *Statement* = an object that can accept a string that is a SQL statement and can execute the string.
 2. *PreparedStatement* = an object that has an associated SQL statement ready to execute.
- Created by methods `createStatement()` (or `prepareStatement(Q)` for prepared statements).

Creating Statements

- The Connection class has methods to create Statements and PreparedStatements.

```
Statement stat1 = myCon.createStatement();
```

```
PreparedStatement stat2 =
```

```
myCon.createStatement(
```

`createStatement` with no argument returns a Statement; with one argument it returns a PreparedStatement.

```
    "SELECT beer, price FROM Sells " +
```

```
    "WHERE bar = 'Joe' 's Bar' "
```

```
);
```

Executing SQL Statements

- JDBC distinguishes queries from modifications, all of which it calls “updates.”
- Statement and PreparedStatement each have methods `executeQuery` and `executeUpdate`.
 - For Statements: one argument, consisting of the query or modification to be executed.
 - For PreparedStatements: no argument, since a prepared statement already has an associated object.

The 4 “Execute” Methods

- `executeQuery(Q)` takes a statement Q, which must be a query, that is applied to a Statement object. Returns a ResultSet object, the multiset of tuples produced by Q.
- `executeQuery()` is applied to a PreparedStatement object.
- `executeUpdate(U)`. Takes a database modification U, and when applied to a Statement object, executes U. No ResultSet is returned
- `executeUpdate()` is applied to a PreparedStatement object.

Example: Update

- stat1 is a Statement.
- We can use it to insert a tuple as:

```
stat1.executeUpdate(  
    "INSERT INTO Sells " +  
    "VALUES('Brass Rail','Export',3.00)"  
);
```

Example: Query

- stat2 is a PreparedStatement holding the query "SELECT beer, price FROM Sells WHERE bar = 'Joe's Bar'".
- `executeQuery` returns an object of class `ResultSet` (next slide)
- The query:

```
ResultSet menu = stat2.executeQuery();
```

Accessing the ResultSet

- An object of type *ResultSet* is something like a *cursor* (which we'll see later).
- Aside: A **cursor** is essentially a tuple-variable that ranges over all tuples in the result of some query.
 - Using a cursor lets one successively iterate through tuples satisfying a query.
- Method **next()** advances the “cursor” to the next tuple.
 - The first time **next()** is applied, it gets the first tuple.
 - If there are no more tuples, **next()** returns the value **false**.

Accessing Components of Tuples

- When a ResultSet is referring to a tuple, we can get the components of that tuple by applying certain methods to the ResultSet.
- Method `getX(i)`, where X is some type, and i is the component number, returns the value of that component.
 - Examples: `getString(i)`, `getInt(i)`, `getFloat(i)`, etc.
 - The value must have type X .

Example: Accessing Components

- menu = ResultSet for query “SELECT beer, price
FROM Sells
WHERE bar = 'Joe's Bar' ”.
- Access beer and price from each tuple by:

```
while ( menu.next() ) {  
    theBeer = menu.getString(1);  
    thePrice = menu.getFloat(2);  
    /* do something with theBeer and thePrice */  
}
```

Accessing Components of Tuples (Method 2)

- Method `getX (ColumnName)`, where X is some type, and *ColumnName* is the component number, returns the value of that component.
 - The value must have type X .

Example: Accessing Components

- menu = ResultSet for query “SELECT beer, price
FROM Sells
WHERE bar = 'Joe's Bar' ”.
- Access beer and price from each tuple by:

```
while ( menu.next() ) {  
    theBeer = Menu.getString(“beer”);  
    thePrice = Menu.getFloat(“price”);  
    /* do something with theBeer and thePrice */  
}
```

SQL Programming: SQL/CLI

- SQL/CLI is a library which provides access to DBMS for C programs.
 - The library for C is called SQL/CLI = “*Call-Level Interface*.”
 - The concepts here are similar to JDBC.

Data Structures

- C connects to the database by records (structs) of the following types:
 1. *Environments* : represent the DBMS installation.
 2. *Connections* : logins to the database.
 3. *Statements* : SQL statements to be passed to a connection.
 4. *Descriptions* : records about tuples from a query, or parameters of a statement.
 - Will ignore here.
- Each of these records is represented by a **handle**, or pointer to the record.
- The header file sqlcli.h provides types for the handles of environments, etc.

Handles

- Function `SQLAllocHandle(T,I,O)` is used to create these structs, which are called environment, connection, and statement **handles**.
 - T = type, e.g., `SQL_HANDLE_STMT`. Also `SQL_HANDLE_ENV`,
`SQL_HANDLE_DBC`
 - I = input handle
= struct at next higher level (statement < connection < environment).
`SQL_NULL_HANDLE` for an environment
E.g. if you want a statement handle, then I is the handle of the “host”
connection
 - O = (address of) output handle (created by `SQLAllocHandle`)

Example: SQLAllocHandle

```
SQLAllocHandle(SQL_HANDLE_STMT, myCon, &myStat);
```

- **myCon** is a previously created connection handle.
- **myStat** is the name of the statement handle that will be created.

Preparing and Executing

- **SQLPrepare**(H , S , L) causes the string S , of length L , to be interpreted as a SQL statement and optimized; the executable statement is placed in statement handle H .
- **SQLExecute**(H) causes the SQL statement represented by statement handle H to be executed.

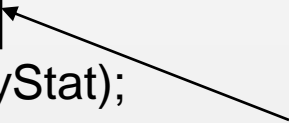
Example: Prepare and Execute

```
SQLPrepare(myStat, "SELECT beer, price FROM Sells
```

```
WHERE bar = 'Joe's Bar',
```

```
SQL_NTS);
```

```
SQLExecute(myStat);
```



This constant says the second argument is a “null-terminated string”; i.e., figure out the length by counting characters.

Direct Execution

- If we are going to execute a statement S only once, we can combine PREPARE and EXECUTE with:

`SQLExecuteDirect(H,S,L);`

- As before, H is a statement handle and L is the length of string S .

Fetching Tuples

- When the SQL statement executed is a query, we need to fetch the tuples of the result.
 - A cursor is implied by the fact we executed a query; the cursor need not be declared.
- **SQLFetch(H)** gets the next tuple from the result of the statement with handle *H*.

Accessing Query Results

- When we fetch a tuple, we need to put the components somewhere.
- Each component is bound to a variable by the function **SQLBindCol**.
- This function has 6 arguments, of which we shall show only 1, 2, and 4:
 - 1 = handle of the query statement.
 - 2 = column number.
 - 4 = address of the variable.

Example: Binding

- Suppose we have just done `SQLExecute(myStat)`, where `myStat` is the handle for query

```
SELECT beer, price
```

```
FROM Sells
```

```
WHERE bar = 'Joe''s Bar'
```

- Bind the result to `theBeer` and `thePrice`:

```
SQLBindCol(myStat, 1, , &theBeer, , );
```

```
SQLBindCol(myStat, 2, , &thePrice, , );
```

Example: Fetching


- Now, we can fetch all the tuples of the answer by:

```
while ( SQLFetch(myStat) != SQL_NO_DATA)
```

```
{
```

```
    /* do something with theBeer and  
       thePrice */
```

```
}
```



CLI macro representing
SQLSTATE = 02000 = “failed
to find a tuple.”

SQL Programming: Embedded SQL

- **Key idea:** A preprocessor turns SQL statements into procedure calls that fit with the surrounding host-language code.
- All embedded SQL statements begin with EXEC SQL, so the preprocessor can find them easily.

Shared Variables

- To connect SQL and the host-language program, the two parts must share some variables.

- Declarations of shared variables are bracketed by:

```
EXEC SQL BEGIN DECLARE SECTION;
```

```
<host-language declarations>
```

```
EXEC SQL END DECLARE SECTION;
```

Use of Shared Variables

- In SQL, the shared variables must be preceded by a colon.
 - They may be used as constants provided by the host-language program.
 - They may get values from SQL statements and pass those values to the host-language program.
- In the host language, shared variables behave like any other variable.
 - Not preceded by a colon here.

Example: Looking Up Prices

- We'll use C with embedded SQL to sketch the important parts of a function that given a beer and a bar, looks up the price of that beer at that bar.
 - Note that a query here returns a single value (tuple)
- Assumes database has our usual **Sells(bar, beer, price)** relation.

Example: C Plus SQL

Single row SELECT Statements:

```
EXEC SQL BEGIN DECLARE SECTION;
```

```
char theBar[21], theBeer[21];
```

```
float thePrice;
```

21-char arrays needed
for 20 chars +
endmarker

```
EXEC SQL END DECLARE SECTION;
```

```
/* obtain values for theBar and theBeer */
```

```
EXEC SQL SELECT price INTO :thePrice  
FROM Sells  
WHERE bar = :theBar AND beer = :theBeer;
```

```
/* do something with thePrice */
```

SELECT-INTO: used
for a single row

Embedded Queries

- Embedded SQL (so far) has a limitation regarding queries:
 - SELECT-INTO for a query guaranteed to produce a single tuple.
 - Otherwise, you have to use a **cursor**.
- Recall: A **cursor** is essentially a tuple-variable that ranges over all tuples in the result of some query.
- Using a cursor lets one iterate through tuples satisfying a query.

Cursor Statements

- Declare a cursor *c* with:

EXEC SQL DECLARE *c* CURSOR FOR <query>;

- Open and close cursor *c* with:

EXEC SQL OPEN CURSOR *c*;

EXEC SQL CLOSE CURSOR *c*;

- The OPEN statement causes the query to be evaluated.
- The CLOSE statement causes the database to delete the temporary relation that holds the result of the query.

- Fetch from *c* by:

EXEC SQL FETCH *c* INTO <variable(s)>;

- Repeated calls to FETCH get successive tuples in the query result.
- Macro NOT FOUND is true if and only if the FETCH fails to find a tuple.

Example Cursor

- From within a host language, want to find the names and cities of customers with more than the variable **amount** dollars in some account.
- Specify the query in SQL and declare a *cursor* for it

EXEC SQL

```
DECLARE c CURSOR FOR
SELECT depositor.customer_name, customer_city
FROM depositor, customer, account
WHERE depositor.customer_name = customer.customer_name
      AND depositor account_number = account.account_number
      AND account.balance > :amount ;
```

Updates Through Cursors

- Can update tuples fetched by cursor by declaring that the cursor is for update

```
DECLARE c CURSOR FOR  
  SELECT *  
  FROM account  
  WHERE branch_name = 'Perryridge'  
  FOR UPDATE
```

- To update tuple at the current location of cursor *c*

```
UPDATE account  
  SET balance = balance + 100  
  WHERE CURRENT OF c
```

Example: Print Joe's Menu

- Write C + SQL to print Joe's menu – the list of beer-price pairs that we find in `Sells(bar, beer, price)` with `bar = Joe's Bar`.
- A cursor will visit each Sells tuple that has `bar = Joe's Bar`.

Example: Declarations

```
EXEC SQL BEGIN DECLARE SECTION;
```


```
    char theBeer[21]; float thePrice;
```

```
EXEC SQL END DECLARE SECTION;
```

```
EXEC SQL DECLARE c CURSOR FOR
```

```
    SELECT beer, price FROM Sells
```

```
    WHERE bar = 'Joe's Bar';
```



The cursor declaration goes
outside the declare-section

Example: Executable Part

```
EXEC SQL OPEN CURSOR c;
while(1) {
    EXEC SQL FETCH c
        INTO :theBeer, :thePrice;
    if (NOT FOUND) break;
    /* format and print theBeer and thePrice */
}
EXEC SQL CLOSE CURSOR c;
```

Need for Dynamic SQL

- Most applications use specific queries and modification statements to interact with the database.
 - The DBMS compiles EXEC SQL ... statements into specific procedure calls and produces an ordinary host-language program that uses a library.
- Dynamic SQL allows programs to construct and submit SQL queries at run time.

Dynamic SQL

- Preparing a query:

EXEC SQL PREPARE <query-name>

FROM <text of the query>;

- Executing a query:

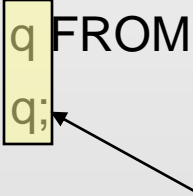
EXEC SQL EXECUTE <query-name>;

- “Prepare” = optimize query.
- Prepare once, execute many times.

Example: A Generic Interface

```
EXEC SQL BEGIN DECLARE SECTION;  
    char query[MAX_LENGTH];  
EXEC SQL END DECLARE SECTION;
```

```
while(1) {  
    /* issue SQL> prompt */  
    /* read user's query into array query */  
    EXEC SQL PREPARE q FROM :query;  
    EXEC SQL EXECUTE q;  
}
```



q is an SQL variable
representing the
optimized form of
whatever statement
is typed into :query

Execute-Immediate

- If we are only going to execute the query once, we can combine the PREPARE and EXECUTE steps into one.

- Use:

EXEC SQL EXECUTE IMMEDIATE <text>;

Example: Generic Interface Again

```
EXEC SQL BEGIN DECLARE SECTION;
```

```
    char query[MAX_LENGTH];
```

```
EXEC SQL END DECLARE SECTION;
```

```
while(1) {
```

```
    /* issue SQL> prompt */
```

```
    /* read user's query into array query */
```

```
    EXEC SQL EXECUTE IMMEDIATE :query;
```

```
}
```

Another Example

Example of the use of dynamic SQL from within a C program.

```
CHAR *sqlprog = "UPDATE account  
                SET balance = balance * 1.05  
                WHERE account_number = ?"  
EXEC SQL PREPARE dynprog FROM :sqlprog;  
CHAR account[10] = "A-101";  
EXEC SQL EXECUTE dynprog USING :account;
```

- The dynamic SQL program contains a ?, which is a place holder for a value that is provided when the SQL program is executed.