



Relational Algebra & Calculus




Relational Query Languages

- ❖ Query languages: Allow **manipulation** and **retrieval of data** from a database.
- ❖ Relational model supports simple, powerful QLs:
 - Strong formal foundation based on logic.
 - Allows for much optimization.
- ❖ Query Languages **!=** programming languages!
 - QLs not expected to be “Turing complete”.
 - QLs not intended to be used for complex calculations.
 - QLs support easy, efficient access to large data sets.



Formal Relational Query Languages


- ❖ Two mathematical Query Languages form the basis for “real” languages (e.g. SQL), and for implementation:
 - Relational Algebra: More **operational** (procedural), very useful for representing execution plans.
 - Relational Calculus: Lets users describe what they want, rather than how to compute it: **Non-operational, declarative**.



Preliminaries

- ❖ A query is applied to *relation instances*, and the result of a query is also a relation instance.
 - *Schemas of input* relations for a query are *fixed*.
 - The *schema for the result* of a given query is also *fixed!* - determined by definition of query language constructs.
- ❖ Positional vs. named-field notation:
 - Positional notation easier for formal definitions, named-field notation more readable.
 - Both used in SQL

Example Instances




<i>R1</i>	<u>sid</u>	<u>bid</u>	<u>day</u>
	22	101	10/10/96
	58	103	11/12/96

- ❖ “Sailors” and “Reserves” relations for our examples.
- ❖ We’ll use positional or named field notation, assume that **names of fields in query results are ‘inherited’ from names of fields in query input relations.**

<i>S1</i>	<u>sid</u>	sname	rating	age
	22	dustin	7	45.0
	31	lubber	8	55.5
	58	rusty	10	35.0

<i>S2</i>	<u>sid</u>	sname	rating	age
	28	yuppy	9	35.0
	31	lubber	8	55.5
	44	guppy	5	35.0
	58	rusty	10	35.0



Relational Algebra

❖ Basic operations:

- Selection (σ) Selects a subset of rows from relation.
- Projection (π) Deletes unwanted columns from relation.
- Cross-product (\times) Allows us to combine two relations.
- Set-difference ($-$) Tuples in reln. 1, but not in reln. 2.
- Union (\cup) Tuples in reln. 1 and in reln. 2.

❖ Additional operations:

- Intersection, join, division, renaming: Not essential, but (very!) useful.

❖ Since each operation returns a relation, **operations can be composed**: algebra is “closed”.

Projection

- ❖ Deletes attributes that are not in *projection list*.
- ❖ *Schema* of result contains exactly the fields in the projection list, with the same names that they had in the input relation.
- ❖ Projection operator has to eliminate *duplicates*! Why?
 - Note: real systems typically don't do duplicate elimination unless the user explicitly asks for it (by DISTINCT). Why not?



sname	rating
yuppy	9
lubber	8
guppy	5
rusty	10

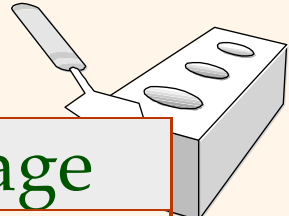
$\pi_{sname, rating}(S2)$

age
35.0
55.5

$\pi_{age}(S2)$

Selection

- ❖ Selects rows that satisfy *selection condition*.
- ❖ No duplicates in result!
Why?
- ❖ *Schema* of result identical to schema of input relation.
- ❖ What is Operator composition?
- ❖ Selection is distributive over binary operators
- ❖ Selection is commutative




sid	sname	rating	age
28	yuppy	9	35.0
58	rusty	10	35.0

$$\sigma_{rating > 8}(S2)$$

sname	rating
yuppy	9
rusty	10

$$\pi_{sname, rating}(\sigma_{rating > 8}(S2))$$

Union, Intersection, Set-Difference



- ❖ All of these operations take two input relations, which must be union-compatible:
 - Same number of fields.
 - ‘Corresponding’ fields have the same type.
- ❖ What is the *schema* of result?

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0
44	guppy	5	35.0
28	yuppy	9	35.0

$S1 \cup S2$


sid	sname	rating	age
22	dustin	7	45.0

$S1 - S2$

sid	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35.0

$S1 \cap S2$


Cross-Product (Cartesian Product)



- ❖ Each row of S1 is paired with each row of R1.
- ❖ *Result schema* has one field per field of S1 and R1, with field names 'inherited' if possible.
 - *Conflict*: Both S1 and R1 have a field called *sid*.

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

- Renaming operator: $\rho (C(1 \rightarrow sid1, 5 \rightarrow sid2), S1 \times R1)$




Joins: used to combine relations

❖ Condition Join: $R \bowtie_c S = \sigma_c (R \times S)$

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	58	103	11/12/96

$$S1 \bowtie_{S1.sid < R1.sid} R1$$

- ❖ *Result schema* same as that of cross-product.
- ❖ Fewer tuples than cross-product, might be able to compute more efficiently
- ❖ Sometimes called a *theta-join*.




Join

- ❖ Equi-Join: A special case of condition join where the condition c contains only *equalities*.

sid	sname	rating	age	bid	day
22	dustin	7	45.0	101	10/10/96
58	rusty	10	35.0	103	11/12/96

$$S1 \bowtie_{sid} R1$$


- ❖ Result schema similar to cross-product, but only one copy of fields for which equality is specified.
- ❖ Natural Join: Equijoin on *all* common fields.




Properties of join

- ❖ Selecting power: can join be used for selection?
- ❖ Is join commutative? $S1 \bowtie R1 = R1 \bowtie S1$?
- ❖ Is join associative? $S1 \bowtie (R1 \bowtie C1) = (S1 \bowtie R1) \bowtie C1$?
- ❖ Join and projection perform complementary functions
- ❖ Lossless and lossy decomposition

Division



- ❖ Not supported as a primitive operator, but useful for expressing queries like:
Find sailors who have reserved all boats.
- ❖ Let A have 2 fields, x and y ; B have only field y :
 - $A/B = \{ \langle x \rangle \mid \exists \langle x, y \rangle \in A \ \forall \langle y \rangle \in B \}$
 - i.e., **A/B contains all x tuples (sailors) such that for every y tuple (boat) in B , there is an xy tuple in A .**
 - Or: If the set of y values (boats) associated with an x value (sailor) in A contains all y values in B , the x value is in A/B .
- ❖ In general, x and y can be any lists of fields; y is the list of fields in B , and $x \cup y$ is the list of fields of A .



Examples of Division A/B

sno	pno
s1	p1
s1	p2
s1	p3
s1	p4
s2	p1
s2	p2
s3	p2
s4	p2
s4	p4

A

pno
p2

$B1$

sno
s1
s2
s3
s4

$A/B1$

pno
p2
p4

$B2$

sno
s1
s4


$A/B2$

pno
p1
p2
p4

$B3$


sno
s1

$A/B3$



Example of Division

- ❖ Find all customers who have an account at all branches located in Chville
 - Branch (bname, assets, bcity)
 - Account (bname, acct#, cname, balance)



Example of Division

R1: Find all branches in Chville


R2: Find (bname, cname) pair from Account

R3: Customers in r2 with every branch name in r1

$$r1 = \pi_{bname}(\sigma_{bcity='Chville'} Branch)$$

$$r2 = \pi_{bname, cname}(Account)$$

$$r3 = r2 \div r1$$




Expressing A/B Using Basic Operators

- ❖ Division is not essential op; just a useful shorthand.
 - Also true of joins, but joins are so common that systems implement joins specially.
- ❖ *Idea*: For A/B , compute all x values that are not 'disqualified' by some y value in B .
 - x value is *disqualified* if by attaching y value from B , we obtain an xy tuple that is not in A .

Disqualified x values: $\pi_x((\pi_x(A) \times B) - A)$

A/B : $\pi_x(A) -$ all disqualified tuples



Exercises


Given relational schema:

Sailors (sid, sname, rating, age)

Reservation (sid, bid, _date)


Boats (bid, bname, color)

- 1) Find names of sailors who've reserved boat #103
- 2) Find names of sailors who've reserved a red boat
- 3) Find sailors who've reserved a red or a green boat
- 4) Find sailors who've reserved a red and a green boat
- 5) Find the names of sailors who've reserved all boats




Summary of Relational Algebra

- ❖ The relational model has rigorously defined query languages that are simple and powerful.
- ❖ Relational algebra is more operational; useful as internal representation for query evaluation plans.
- ❖ Several ways of expressing a given query; a query optimizer should choose the most efficient version.



Relational Calculus

- ❖ Comes in two flavors: Tuple relational calculus (TRC) and Domain relational calculus (DRC).
- ❖ Calculus has *variables*, *constants*, *comparison ops*, *logical connectives* and *quantifiers*.
 - TRC: Variables range over (i.e., get bound to) *tuples*.
 - DRC: Variables range over *domain elements* (= field values).
 - Both TRC and DRC are simple subsets of first-order logic.
- ❖ Expressions in the calculus are called *formulas*. An answer tuple is essentially an assignment of constants to variables that make the formula evaluate to *true*.




Domain Relational Calculus

- ❖ *Query* has the form:

$$\left\{ \langle x_1, x_2, \dots, x_n \rangle \mid p(\langle x_1, x_2, \dots, x_n \rangle) \right\}$$

- ❖ *Answer* includes all tuples $\langle x_1, x_2, \dots, x_n \rangle$ that make the formula $p(\langle x_1, x_2, \dots, x_n \rangle)$ be true.
- ❖ *Formula* is recursively defined, starting with simple *atomic formulas* (getting tuples from relations or making comparisons of values), and building bigger and better formulas using the *logical connectives*.




DRC Formulas

❖ Atomic formula:

- $\langle x_1, x_2, \dots, x_n \rangle \in Rname$, or $X \text{ op } Y$, or $X \text{ op } \text{constant}$
- *op* is one of $<, >, =, \leq, \geq, \neq$

❖ Formula:

- an atomic formula, or
 - $\neg p, p \wedge q, p \vee q$, where p and q are formulas, or
 - $\exists X (p(X))$, where X is *a domain variable* or
 - $\forall X (p(X))$, where X is *a domain variable*.
- ❖ The use of **quantifiers** $\exists X$ and $\forall X$ is said to **bind** X .




Free and Bound Variables

- ❖ The use of **quantifiers** $\exists X$ and $\forall X$ in a formula is said to bind X .
 - A variable that is **not bound** is free.
- ❖ Let us revisit the definition of a **query**:

$$\left\{ \langle x_1, x_2, \dots, x_n \rangle \mid p(\langle x_1, x_2, \dots, x_n \rangle) \right\}$$

- ❖ There is an important restriction: the variables x_1, \dots, x_n that appear to the left of `|' must be the *only* free variables in the formula $p(\dots)$.


Find all sailors with a rating above 7



$$\left\{ \langle I, N, T, A \rangle \mid \langle I, N, T, A \rangle \in \text{Sailors} \wedge T > 7 \right\}$$

- ❖ The condition $\langle I, N, T, A \rangle \in \text{Sailors}$ ensures that the domain variables I , N , T and A are bound to fields of the same Sailors tuple.
- ❖ The term $\langle I, N, T, A \rangle$ to the left of `|` (which should be read as *such that*) says that every tuple $\langle I, N, T, A \rangle$ that satisfies $T > 7$ is in the answer.
- ❖ Modify this query to answer:
 - Find sailors who are older than 18 or have a rating under 9, and are called 'Joe'.


*Find sailors rated > 7 who have reserved
boat #103*



$$\left\{ \langle I, N, T, A \rangle \mid \langle I, N, T, A \rangle \in \text{Sailors} \wedge T > 7 \wedge \right. \\ \left. \exists Ir, Br, D \left(\langle Ir, Br, D \rangle \in \text{Reserves} \wedge Ir = I \wedge Br = 103 \right) \right\}$$

- ❖ We have used $\exists Ir, Br, D \dots$ as a shorthand for $\exists Ir \left(\exists Br \left(\exists D \dots \right) \right)$
- ❖ Note the use of \exists to find a tuple in Reserves that ‘joins with’ the Sailors tuple under consideration.


Find sailors rated > 7 who've reserved a red boat



$$\left\{ \left\langle I, N, T, A \right\rangle \mid \left\langle I, N, T, A \right\rangle \in \text{Sailors} \wedge T > 7 \wedge \right. \\ \left. \exists Ir, Br, D \left(\left\langle Ir, Br, D \right\rangle \in \text{Reserves} \wedge Ir = I \wedge \right. \right. \\ \left. \left. \exists B, BN, C \left(\left\langle B, BN, C \right\rangle \in \text{Boats} \wedge B = Br \wedge C = 'red' \right) \right) \right\}$$

- ❖ Observe how the parentheses control the scope of each quantifier's binding.
- ❖ This may look cumbersome, but with a good user interface, it could be intuitive. (MS Access, QBE)


Find sailors who've reserved all boats



$$\left\{ \langle I, N, T, A \rangle \mid \langle I, N, T, A \rangle \in \text{Sailors} \wedge \right. \\ \left. \forall B, BN, C \left(\neg \left(\langle B, BN, C \rangle \in \text{Boats} \right) \vee \right. \right. \\ \left. \left. \left(\exists Ir, Br, D \left(\langle Ir, Br, D \rangle \in \text{Reserves} \wedge I = Ir \wedge Br = B \right) \right) \right) \right\}$$

- ❖ Find all sailors I such that for each 3-tuple $\langle B, BN, C \rangle$ either it is not a tuple in Boats or there is a tuple in Reserves showing that sailor I has reserved it.

Find sailors who've reserved all boats (again!)



$$\left\{ \langle I, N, T, A \rangle \mid \langle I, N, T, A \rangle \in \text{Sailors} \wedge \right. \\ \left. \forall \langle B, BN, C \rangle \in \text{Boats} \right. \\ \left. \left(\exists \langle Ir, Br, D \rangle \in \text{Reserves} (I = Ir \wedge Br = B) \right) \right\}$$


❖ Simpler notation, same query. (Much clearer!)

❖ To find sailors who've reserved all red boats:


$$\dots \left(C \neq 'red' \vee \exists \langle Ir, Br, D \rangle \in \text{Reserves} (I = Ir \wedge Br = B) \right)$$

Any other way to specify it? Equivalence in logic

Unsafe Queries, Expressive Power



- ❖ It is possible to write syntactically correct calculus queries that have an infinite number of answers! Such queries are called unsafe.
 - e.g., $\{S \mid \neg (S \in Sailors)\}$
- ❖ It is known that every query that can be expressed in relational algebra can be expressed as a safe query in DRC / TRC; the converse is also true.
- ❖ Relational Completeness: Query language (e.g., SQL) can express every query that is expressible in relational algebra/calculus.



Exercise of tuple calculus


Given relational schema:

Sailors (sid, sname, rating, age)

Reservation (sid, bid, date)

Boats (bid, bname, color)

- 1) Find all sailors with a rating above 7.
- 2) Find the names and ages of sailors with a rating above 7
- 3) Find the sailor name, boat id, and reservation date for each reservation
- 4) Find the names of the sailors who reserved all boats.



Summary of Relational Calculus

- ❖ Relational calculus is non-operational, and users define queries in terms of what they want, not in terms of how to compute it. (Declarativeness.)
- ❖ Algebra and safe calculus have same expressive power, leading to the notion of relational completeness.