

Relational Algebra



SQL

Relational Query Languages

- Languages for describing queries on a relational database
- *Structured Query Language* (SQL)
 - Predominant application-level query language
 - Declarative
- *Relational Algebra*
 - Intermediate language used within DBMS
 - Procedural

SQL

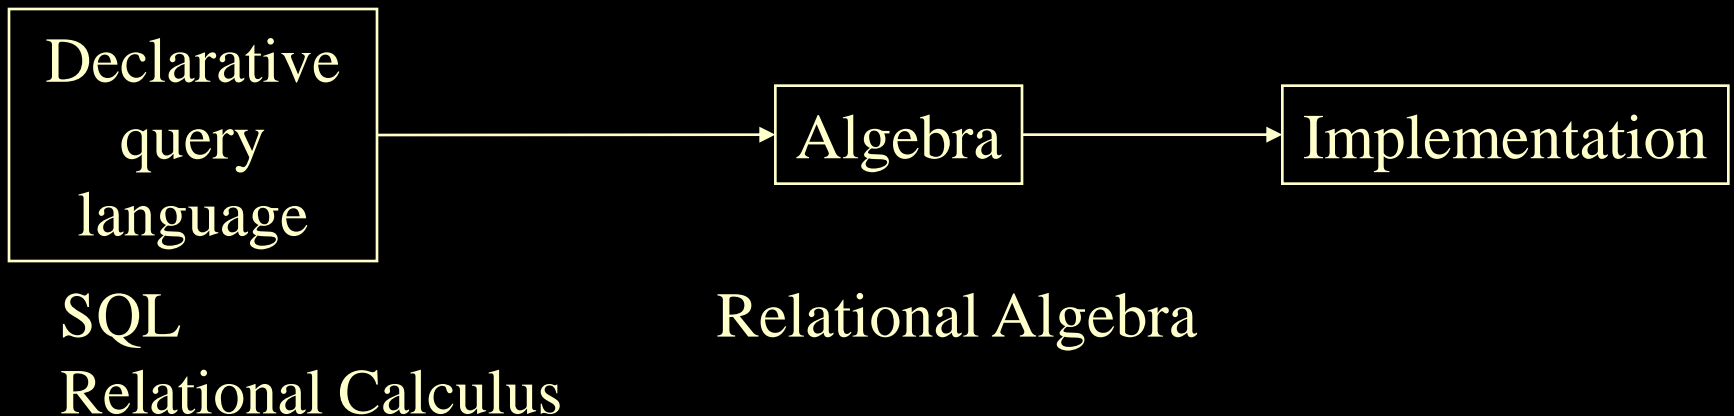


SQL consists of the following parts:

- Data Definition Language (DDL)
- Interactive Data Manipulation Language (Interactive DML)
- Embedded Data Manipulation Language (Embedded DML)
- Views
- Integrity
- Transaction Control
- Authorization
- Catalog and Dictionary Facilities

Relational Algebra

- Formalism for creating new relations from existing ones
- Its place in the big picture:



Relational Algebra



- Five operators:
 - Union: \cup
 - Difference: $-$
 - Selection: σ
 - Projection: Π
 - Cartesian Product: \times
- Derived or auxiliary operators:
 - Intersection, complement
 - Joins (natural, equi-join, etc.)
 - Renaming: ρ

Renaming

- Changes the schema, not the instance
- Notation: $\rho_{B_1, \dots, B_n}(R)$
- Example:
 - $\rho_{\text{LastName}, \text{SocSocNo}}(\text{Employee})$
 - Output schema:
(LastName, SocSocNo)

Renaming Example

Employee

Name	SSN
John	999999999
Tony	777777777

$\rho_{\text{LastName, SocSocNo}}$ (**Employee**)

LastName	SocSocNo
John	999999999
Tony	777777777

Set Operators

- SQL provides UNION, EXCEPT (set difference), and INTERSECT for union compatible tables
- Example: Find all professors in the IT Department and all professors that have taught IT courses

```
(SELECT  P.Name
FROM    Professor P, Teaching T
WHERE   P.Id=T.ProfId AND T.CrsCode LIKE 'IT%')
UNION
(SELECT  P.Name
FROM    Professor P
WHERE   P.DeptId = 'IT')
```


Division

- *Strategy for implementing division in SQL:*
 - Find set, A, of all departments in which a particular professor, p , has taught a course
 - Find set, B, of all departments
 - Output p if $A \supseteq B$, or, equivalently, if $B-A$ is empty

Division – SQL Solution

```
SELECT P.Id
FROM Professor P
WHERE NOT EXISTS
    (SELECT D.DeptId           -- set B of all dept Ids
     FROM Department D
      EXCEPT
     SELECT C.DeptId           -- set A of dept Ids of depts in
                                -- which P has taught a course
     FROM Teaching T, Course C
    WHERE T.ProfId=P.Id       -- global variable
          AND T.CrsCode=C.CrsCode)
```

Aggregates

- Functions that operate on sets:
 - COUNT, SUM, AVG, MAX, MIN
- Produce numbers (not tables)
- Not part of relational algebra

```
SELECT COUNT(*)  
FROM Professor P
```

```
SELECT MAX (Salary)  
FROM Employee E
```

Aggregates

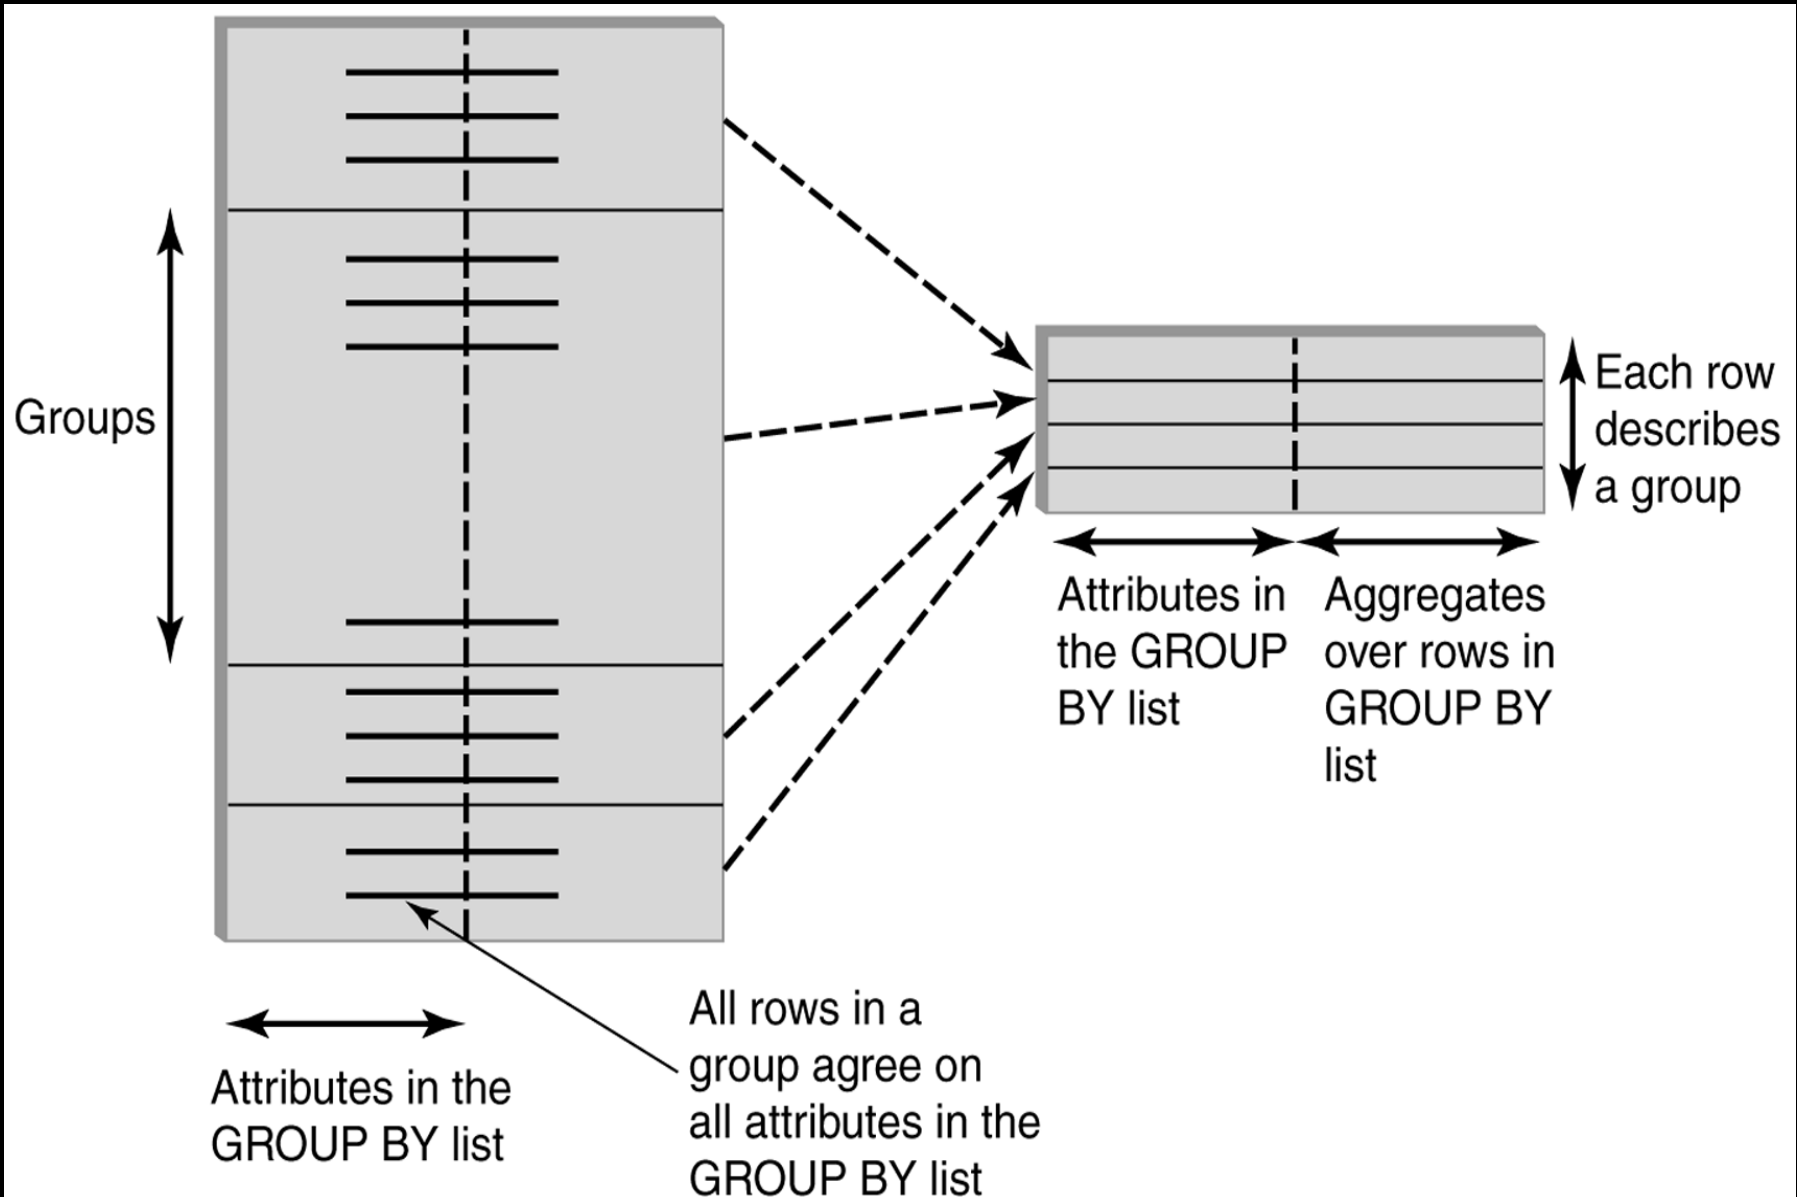
Count the number of courses taught in S2000

```
SELECT COUNT (T.CrsCode)
FROM Teaching T
WHERE T.Semester = 'S2000'
```

But if multiple sections of same course
are taught, use:

```
SELECT COUNT (DISTINCT T.CrsCode)
FROM Teaching T
WHERE T.Semester = 'S2000'
```

GROUP BY

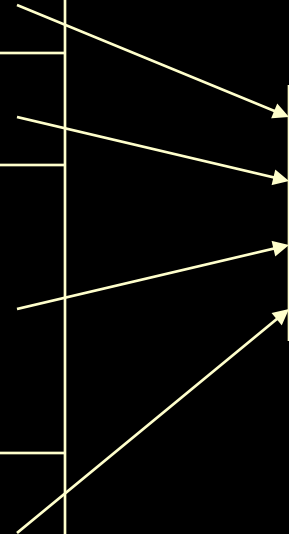


GROUP BY - Example

Transcript

1234	
1234	
1234	
1234	

1234	3.3	4



Attributes:

- student's *Id*
- avg grade
- number of courses

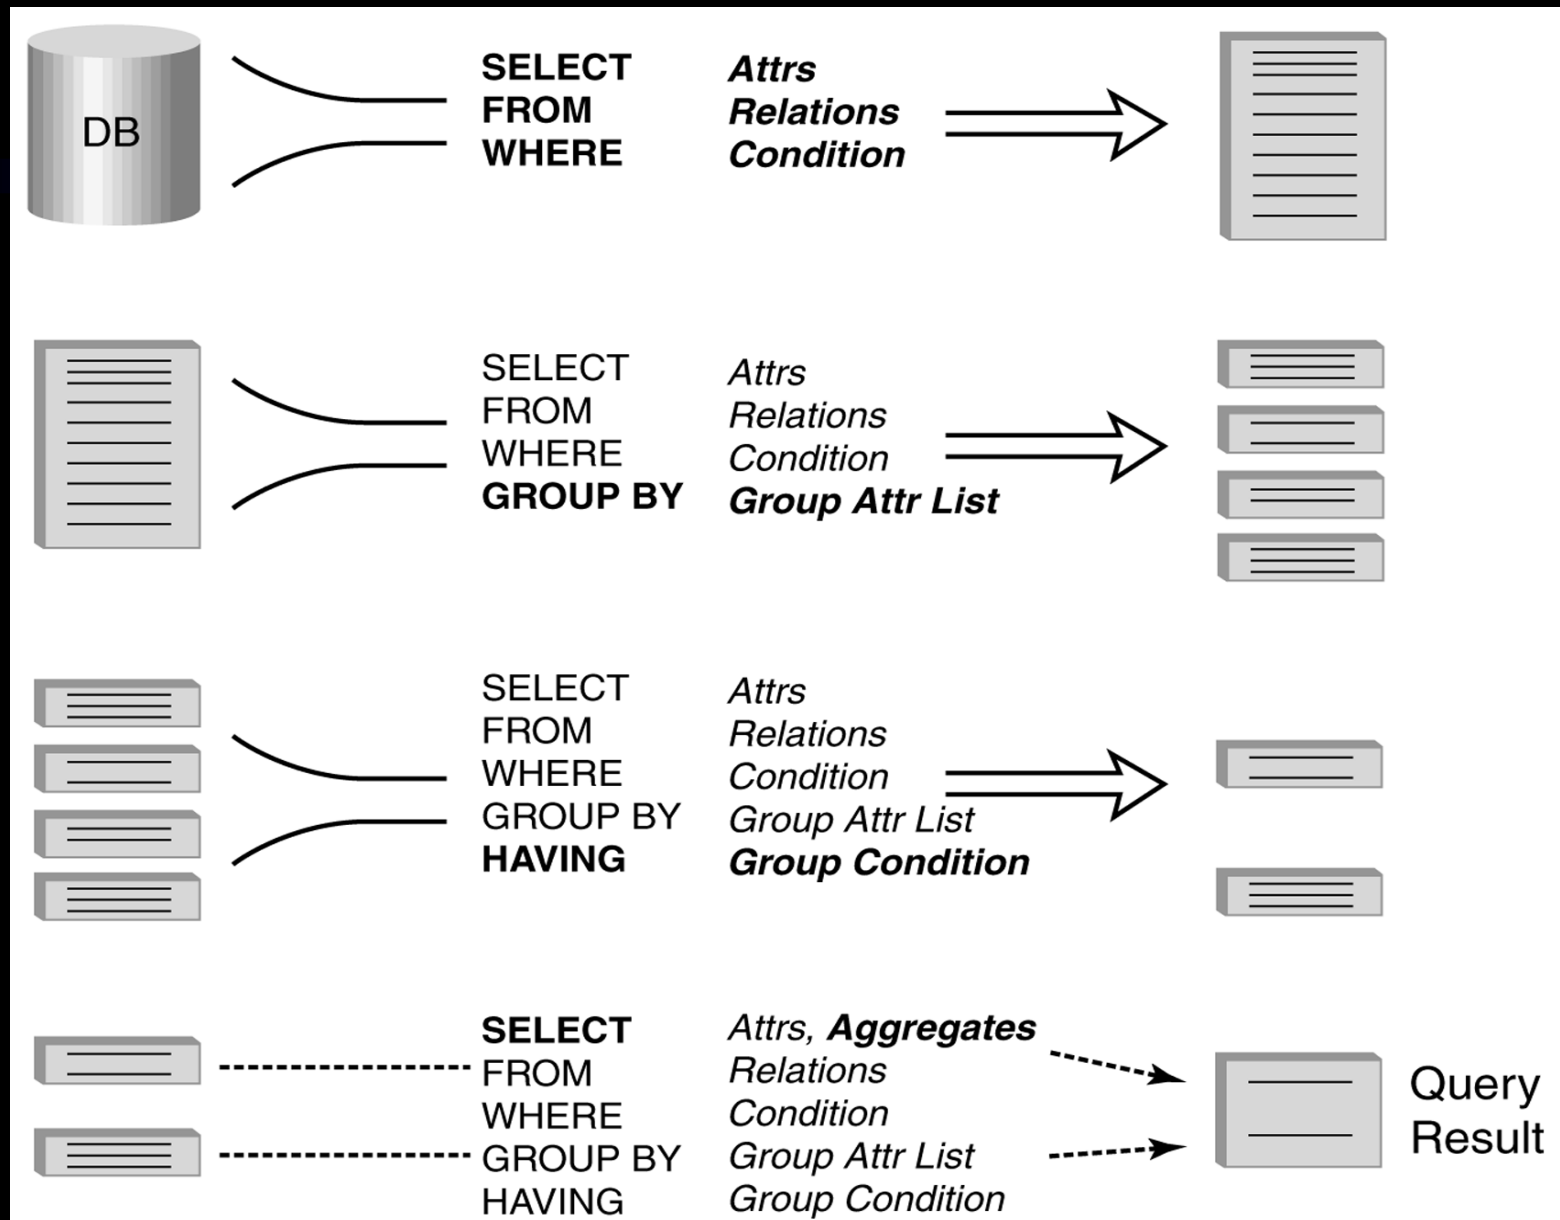
```
SELECT T.StudId, AVG(T.Grade), COUNT (*)  
FROM Transcript T  
GROUP BY T.StudId
```

HAVING Clause

- Eliminates unwanted groups (analogous to WHERE clause)
- HAVING condition constructed from attributes of GROUP BY list and aggregates of attributes not in list

```
SELECT  T.StudId, AVG(T.Grade) AS CumGpa,  
        COUNT (*) AS NumCrS  
FROM    Transcript T  
WHERE   T.CrsCode LIKE 'CS%'  
GROUP BY T.StudId  
HAVING  AVG (T.Grade) > 3.5
```

Evaluation of GroupBy with Having



Example

- Output the name and address of all seniors on the Dean's List

```
SELECT S.Name, S.Address
FROM Student S, Transcript T
WHERE S.StudId = T.StudId AND S.Status = 'senior'

GROUP BY S.StudId -- wrong
         S.Name, S.Address -- right

HAVING AVG (T.Grade) > 3.5 AND SUM (T.Credit) > 90
```

ORDER BY Clause

- Causes rows to be output in a specified order

```
SELECT  T.StudId, COUNT (*) AS NumCrs,  
        AVG(T.Grade) AS CumGpa  
FROM    Transcript T  
WHERE   T.CrsCode LIKE 'CS%'  
GROUP BY T.StudId  
HAVING  AVG (T.Grade) > 3.5  
ORDER BY DESC CumGpa, ASC StudId
```

Views

- Used as a relation, but rows are not physically stored.
 - The contents of a view is *computed* when it is used within an SQL statement
- View is the result of a SELECT statement over other views and base relations
- When used in an SQL statement, the view definition is substituted for the view name in the statement
 - SELECT statement can be nested in FROM clause

View - Example

```
CREATE VIEW CumGpa (StudId, Cum) AS  
  SELECT T.StudId, AVG (T.Grade)  
  FROM Transcript T  
  GROUP BY T.StudId
```

```
SELECT S.Name, C.Cum  
FROM CumGpa C, Student S  
WHERE C.StudId = S.StudId AND C.Cum > 3.5
```

View Benefits

- *Access Control*: Users not granted access to base tables. Instead they are granted access to the view of the database appropriate to their needs.
 - *External schema* is composed of views.
 - View allows owner to provide SELECT access to a subset of columns (analogous to providing UPDATE and INSERT access to a subset of columns)

Modifying Data - Update

```
UPDATE Employee E
SET      E.Salary = E.Salary * 1.05
WHERE    E.Department = 'research'
```

- Updates rows in a single table
- All rows satisfying WHERE clause (general form, including subqueries, allowed) are updated

Updating Views

- Question: Since views look like tables to users, can they be updated?
- Answer: Yes – a view update changes the underlying base table to produce the requested change to the view

```
CREATE VIEW  CsReg (StudId, CrsCode, Semester) AS  
SELECT      T.StudId, T. CrsCode, T.Semester  
FROM        Transcript T  
WHERE       T.CrsCode LIKE 'CS%' AND T.Semester='S2000'
```

Updating Views - Problem 1

INSERT INTO **CsReg** (*StudId, CrsCode, Semester*)
VALUES (1111, 'CSE305', 'S2000')

- **Question:** What value should be placed in attributes of underlying table that have been projected out (e.g., *Grade*)?
- **Answer:** NULL (assuming null allowed in the missing attribute) or DEFAULT

Updating Views - Problem 2

```
INSERT INTO CsReg (StudId, CrsCode, Semester)  
VALUES (1111, 'ECO105', 'S2000')
```

- **Problem:** New tuple not in view
- **Solution:** Allow insertion (assuming the `WITH CHECK OPTION` clause has not been appended to the `CREATE VIEW` statement)

Updating Views - Problem 3

- Update to the view might not *uniquely* specify the change to the base table(s) that results in the desired modification of the view

```
CREATE VIEW ProfDept (PrName, DeName) AS  
SELECT P.Name, D.Name  
FROM Professor P, Department D  
WHERE P.DeptId = D.DeptId
```

Updating Views - Problem 3 (con't)

- Tuple $\langle \text{Smith}, \text{CS} \rangle$ can be deleted from ProfDept by:
 - Deleting row for Smith from Professor (but this is inappropriate if he is still at the University)
 - Deleting row for CS from Department (not what is intended)
 - Updating row for Smith in Professor by setting *DeptId* to null (seems like a good idea)

Updating Views - Restrictions

- Updatable views are restricted to those in which
 - No Cartesian product in FROM clause
 - no aggregates, GROUP BY, HAVING

For example, if we allowed:

```
CREATE VIEW AvgSalary (DeptId, Avg_Sal ) AS
  SELECT  E.DeptId, AVG(E.Salary)
  FROM    Employee E
  GROUP BY E.DeptId
```

then how do we handle:

```
UPDATE AvgSalary
  SET Avg_Sal = 1.1 * Avg_Sal
```

SQL: Join operation

- A join can be specified in the FROM clause which list the two input relations and the WHERE clause which lists the join condition.

Emp		Dept	
ID	State	ID	Division
1000	CA	1001	IT
1001	MA	1002	Sales
1002	TN	1003	Biotech

SQL: Join operation (cont.)

- inner join = join

SELECT *

FROM emp join dept (or FROM emp, dept)

on emp.id = dept.id;

Emp.ID	Emp.State	Dept.ID	Dept.Division
1001	MA	1001	IT
1002	TN	1002	Sales

SQL: Join operation (cont.)

- left outer join = left join

SELECT *

FROM emp left join dept

on emp.id = dept.id;

Emp.ID	Emp.State	Dept.ID	Dept.Division
1000	CA	null	null
1001	MA	1001	IT
1002	TN	1002	Sales

SQL: Join operation (cont.)

- right outer join = right join

SELECT *

FROM emp right join dept

on emp.id = dept.id;

Emp.ID	Emp.State	Dept.ID	Dept.Division
1001	MA	1001	IT
1002	TN	1002	Sales
null	null	1003	Biotech

SQL: Like operation

Pattern matching selection

- % (arbitrary string)

```
SELECT *
```

```
FROM emp
```

```
WHERE ID like '%01';
```

- finds ID that ends with 01, e.g. 1001, 2001, etc

- _ (a single character)

```
SELECT *
```

```
FROM emp
```

```
WHERE ID like '_01_';
```

- finds ID that has the second and third character as 01, e.g. 1010, 1011, 1012, 1013, etc