

# SQL in the Real World

# Interactive vs. Non-Interactive SQL

- *Interactive SQL*: SQL statements input from terminal; DBMS outputs to screen
  - Inadequate for most applications
    - It may be necessary to process the data before output
    - Amount of data returned not known in advance
    - SQL has very limited expressive power (not Turing-complete)
- *Non-interactive SQL*: SQL statements are included in an application program written in a host language, like C, Java, COBOL

# Application Program

- *Host language*: A conventional language (*e.g.*, C, Java) that supplies control structures, computational capabilities, interaction with physical devices
- *SQL*: supplies ability to interact with database.
- *Using the facilities of both*: the application program can act as an intermediary between an unsophisticated user at a terminal and the DBMS

# Preparation

- Before an SQL statement is executed, it must be *prepared*:
  - What indices can be used?
  - In what order should tables be accessed?
  - What constraints should be checked?
- Decisions based on schema, table sizes, etc.
- Result is a *query execution plan*
- Preparation is a complex activity, usually done at run time, justified by the complexity of query processing

# Introducing SQL Into the Application

- SQL statements can be incorporated into an application program in two different ways:
  - *Statement Level Interface* (SLI): Application program is a mixture of host language statements and SQL statements and directives
  - *Call Level Interface* (CLI): Application program is written entirely in host language
    - SQL statements are values of string variables that are passed as arguments to host language (library) procedures

# Statement Level Interface

- SQL statements and directives in the application have a *special syntax* that sets them off from host language constructs
  - e.g., EXEC SQL *SQL\_statement*
- *Precompiler* scans program and translates SQL statements into calls to host language library procedures that communicate with DBMS
- *Host language compiler* then compiles program

# Statement Level Interface

- SQL constructs in an application take two forms:
  - Standard SQL statements (*static* or *embedded* SQL):  
Useful when SQL portion of program is known at compile time
  - Directives (*dynamic* SQL): Useful when SQL portion of program not known at compile time. Application constructs SQL statements *at run time* as values of host language variables that are manipulated by directives
- Precompiler translates statements and directives into arguments of calls to library procedures.

# Call Level Interface

- Application program written entirely in host language (no precompiler)
  - Examples: JDBC, ODBC
- SQL statements are values of string variables constructed *at run time* using host language
  - Similar to dynamic SQL
- Application uses string variables as arguments of library routines that communicate with DBMS
  - e.g. `executeQuery("SQL query statement")`



# Static SQL

```
EXEC SQL BEGIN DECLARE SECTION;
```

```
    unsigned long num_enrolled;
```

```
    char crs_code;
```

```
    char SQLSTATE [6];
```

```
EXEC SQL END DECLARE SECTION;
```

```
.....
```

```
EXEC SQL SELECT C.NumEnrolled
```

```
    INTO :num_enrolled
```

```
    FROM Course C
```

```
    WHERE C.CrsCode = :crs_code;
```

*Variables  
shared by host  
and SQL*

- Declaration section for host/SQL communication
- Colon convention for value (WHERE) and result (INTO) parameters

# Status

```
EXEC SQL SELECT C.NumEnrolled  
    INTO :num_enrolled  
    FROM Course C  
    WHERE C.CrsCode = :crs_code;  
if ( !strcmp (SQLSTATE, "00000") ) {  
    printf ( "statement failed" )  
};
```

# Connections

- To connect to an SQL database, use a connect statement

**CONNECT** TO *database\_name* AS  
*connection\_name* USING *user\_id*

# Transactions

- No explicit statement is needed to begin a transaction
  - A transaction is initiated when the first SQL statement that accesses the database is executed
- The mode of transaction execution can be set with  
`SET TRANSACTION READ ONLY`  
`ISOLATION LEVEL SERIALIZABLE`
- Transactions are terminated with `COMMIT` or `ROLLBACK` statements

# Example: Course Deregistration

```
EXEC SQL CONNECT TO :dbserver;
if ( ! strcmp (SQLSTATE, "000000") ) exit (1);

.....
EXEC SQL DELETE FROM Transcript T
    WHERE T.StudId = :studid AND T.Semester = 'S2000'
        AND T.CrsCode = :crscode;
if (! strcmp (SQLSTATE, "000000") ) EXEC SQL ROLLBACK;
else {
    EXEC SQL UPDATE Course C
        SET C.Numenrolled = C.Numenrolled - 1
        WHERE C.CrsCode = :crscode;
    if (! strcmp (SQLSTATE, "000000") ) EXEC SQL ROLLBACK;
    else EXEC SQL COMMIT;
}
```

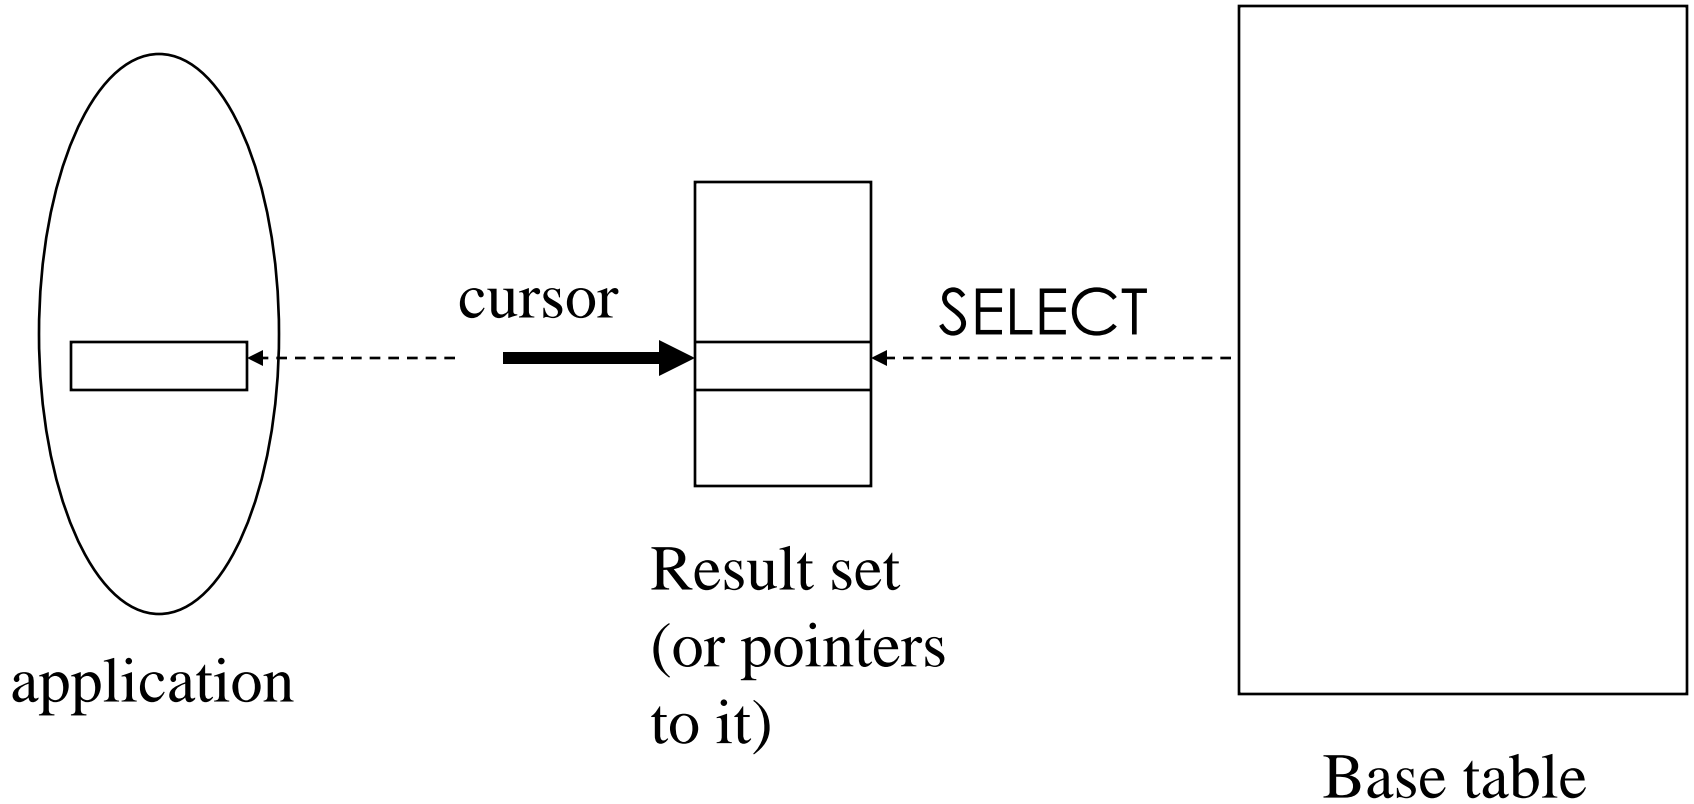
# Buffer Mismatch

- **Problem:** SQL deals with tables (of arbitrary size); host language program deals with fixed size buffers
  - How is the application to allocate storage for the result of a SELECT statement?
- **Solution:** Fetch a single row at a time
  - Space for a single row (number and type of out parameters) can be determined from schema and allocated in application

# Cursor Overview

- *Result set* – set of rows produced by a SELECT statement
- *Cursor* – pointer to a row in the result set.
- Cursor operations:
  - *Declaration*
  - *Open* – execute SELECT to determine result set and initialize pointer
  - *Fetch* – advance pointer and retrieve next row
  - *Close* – deallocate cursor

# Cursor





# Cursor

```
EXEC SQL DECLARE GetEnroll INSENSITIVE CURSOR FOR
  SELECT T.StudId, T.Grade      --cursor is not a schema element
  FROM Transcript T
  WHERE T.CrsCode = :crscode AND T.Semester = 'S2000';
.....
EXEC SQL OPEN GetEnroll;      Value substituted at OPEN time
if ( !strcmp ( SQLSTATE, "00000" )) { ... fail exit... };
.....
EXEC SQL FETCH GetEnroll INTO :studid, :grade;
while ( SQLSTATE = "00000" ) {
  ... process the returned row...
  EXEC SQL FETCH GetEnroll INTO :studid, :grade; }
if ( !strcmp ( SQLSTATE, "02000" )) { ... fail exit... };
.....
EXEC SQL CLOSE GetEnroll;
```

*Reference resolved by compile time,*

# Cursor Types

- *Insensitive cursor*: Result set (effectively) computed and stored in a separate table at OPEN time
  - Changes made to base table subsequent to OPEN (by any transaction) do not affect result set
  - Cursor is read-only
- *Cursors that are not insensitive*: Specification not part of SQL standard
  - Changes made to base table subsequent to OPEN (by any transaction) can affect result set
  - Cursor is updatable

# Insensitive Cursor

cursor



key1	t t t t t t t t
key3	y y y y y y y y
key4	z z z z z z z z

*Result Set*


key1	t t t t t t t t
key2	x x x x x x x x
key3	y y y y y y y y
key4	z z z z z z z z
key5	u u u u u u u u
key6	v v v v v v v v

*Base Table*

# Keyset-Driven Cursor

- Example of a cursor that is not insensitive
- Primary key of each row in result set is computed at open time
- UPDATE or DELETE of a row in base table by a concurrent transaction between OPEN and FETCH might be seen through cursor
- INSERT into base table, however, not seen through cursor
- Cursor is updatable

# Keyset-Driven Cursor



# Cursors

```
DECLARE cursor-name [INSENSITIVE] [SCROLL]
    CURSOR FOR table-expr
    [ ORDER BY column-list ]
    [ FOR {READ ONLY | UPDATE [ OF column-list ] } ]
```

For updatable (not insensitive, not read-only) cursors

```
UPDATE table-name                --base table
```

```
    SET assignment
```

```
    WHERE CURRENT OF cursor-name
```

```
DELETE FROM table-name          --base table
```

```
    WHERE CURRENT OF cursor-name
```

Restriction – *table-expr* must satisfy restrictions of updatable view

# Scrolling

- If SCROLL option not specified in cursor declaration, FETCH always moves cursor forward one position
- If SCROLL option is included in DECLARE CURSOR section, cursor can be moved in arbitrary ways around result set:



*Get previous tuple*

FETCH **PRIOR** FROM GetEnroll INTO :studid, :grade;

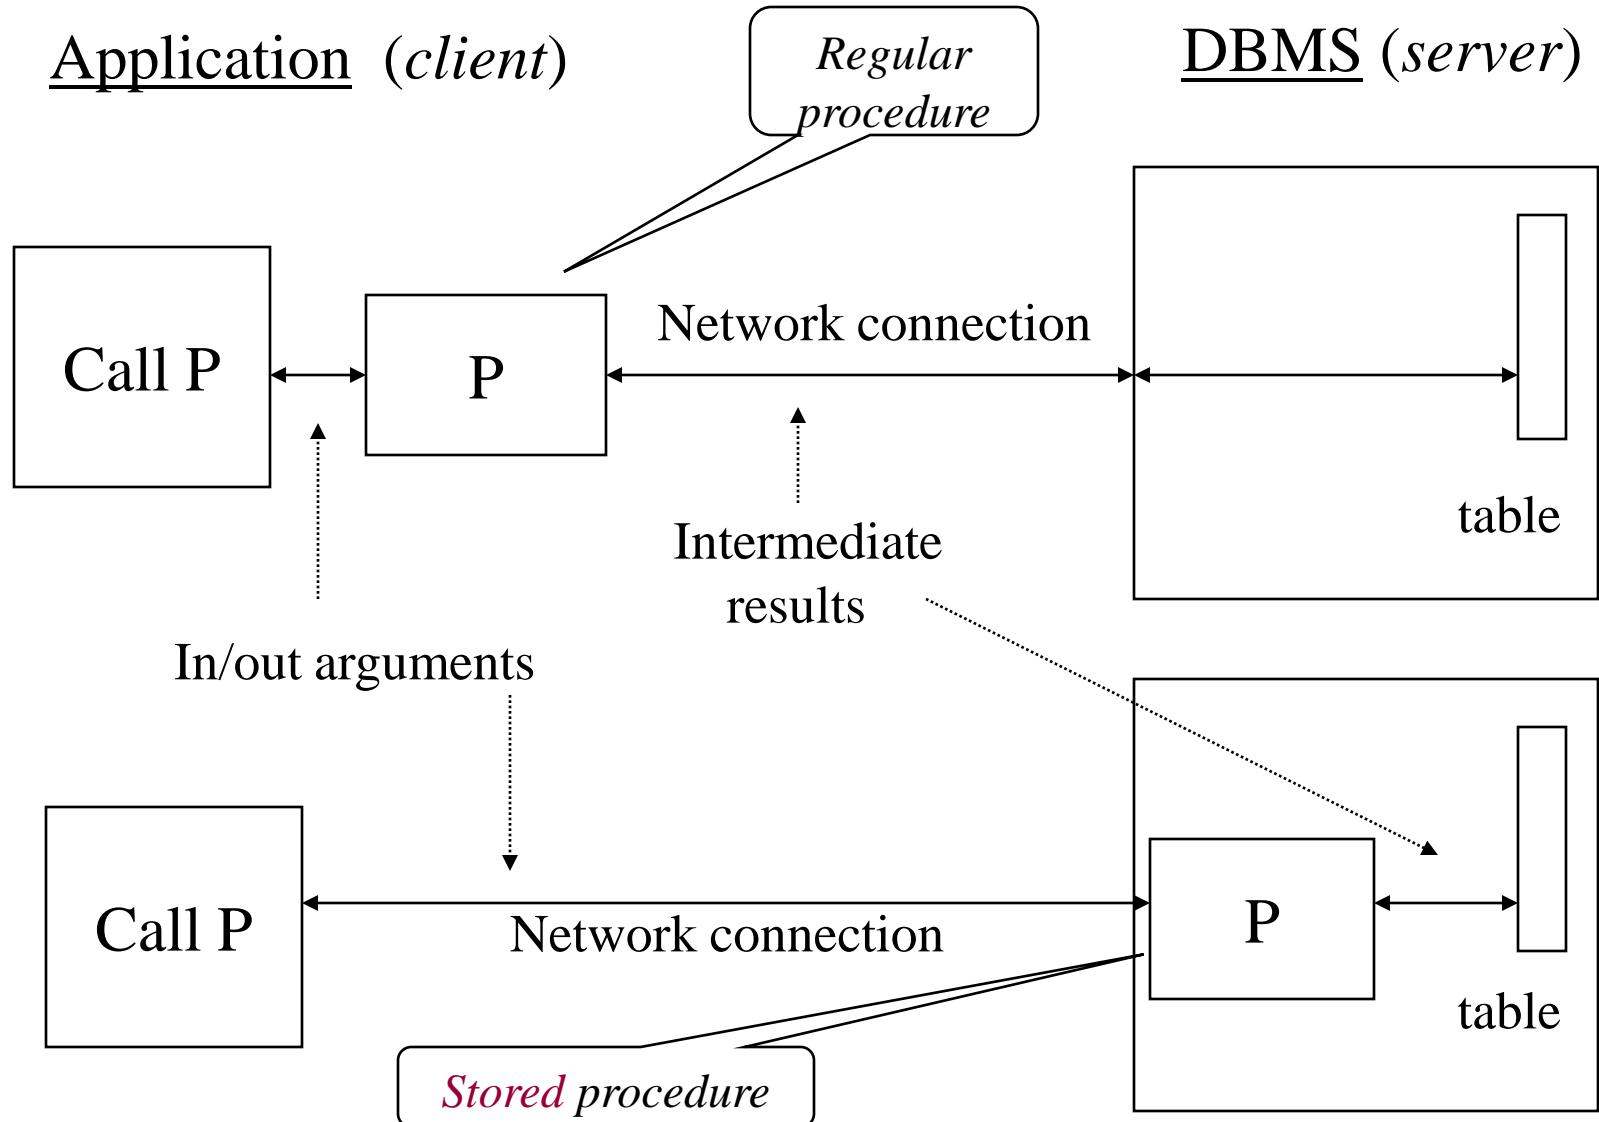
- Also: FIRST, LAST, ABSOLUTE n, RELATIVE n

# Stored Procedures

- Procedure – written in a conventional algorithmic language
  - Included as schema element (stored in DBMS)
  - Invoked by the application
- Advantages:
  - Intermediate data need not be communicated to application (time and cost savings)
  - Procedure's SQL statements prepared in advance
  - Authorization can be done at procedure level
  - Added security since procedure resides in server
  - Applications that call the procedure need not know the details of database schema – all database access is encapsulated within the procedure



# Stored Procedures



# Stored Procedures

*Schema:*

```
CREATE PROCEDURE Register (char :par1, char :par2)
AS BEGIN
    EXEC SQL SELECT ..... ;
    IF ( ..... ) THEN .....      -- SQL embedded in
                                   -- Persistent Stored Modules
                                   -- (PSM) language
    ELSE ....
END
```

*Application:*

```
EXEC SQL EXECUTE PROCEDURE Register ( :crscode, :studid);
```

# Integrity Constraint Checking

- Transaction moves database from an initial to a final state, both of which satisfy all integrity constraints but ...
  - Constraints might not be true of intermediate states hence ...
  - Constraint checks at statement boundaries might be inappropriate
- SQL (optionally) allows checking to be deferred to transaction COMMIT

# Deferred Constraint Checking

*Schema:*

```
CREATE ASSERTION NumberEnrolled  
CHECK ( .....)  
DEFERRABLE;
```

*Application:*

```
SET CONSTRAINT NumberEnrolled DEFERRED;
```

Transaction is aborted if constraint is false at commit time

# Dynamic SQL

- **Problem:** Application might not know in advance:
  - The SQL statement to be executed
  - The database schema to which the statement is directed
- **Example:** User inputs database name and SQL statement interactively from terminal
- In general, application constructs (as the value of a host language string variable) the SQL statement at run time
- Preparation (necessarily) done at run time

# Dynamic SQL

- SQL-92 defines syntax for embedding directives into application for constructing, preparing, and executing an SQL statement
  - Referred to as *Dynamic SQL*
  - Statement level interface
- Dynamic and static SQL can be mixed in a single application

# Dynamic SQL

```
strcpy (tmp, "SELECT C.NumEnrolled FROM Course C  
WHERE C.CrsCode = ?");
```

```
EXEC SQL PREPARE st FROM :tmp;
```

```
EXEC SQL EXECUTE st INTO :num_enrolled USING :crs_code;
```

*placeholder*

- st is an SQL variable; names the SQL statement
- tmp, crscode, num\_enrolled are host language variables (note colon notation)
- crscode is an *in* parameter; supplies value for placeholder (?)
- num\_enrolled is an *out* parameter; receives value from *C.NumEnrolled*

# Dynamic SQL

- PREPARE names SQL statement st and sends it to DBMS for preparation
- EXECUTE causes the statement named st to be executed



# Parameters: Static vs Dynamic SQL

- Static SQL -
  - Names of (host language) parameters are contained in SQL statement and available to precompiler
  - Address and type information in symbol table
  - Routines for fetching and storing argument values can be generated
  - Complete statement (with parameter values) sent to DBMS when statement is executed

```
EXEC SQL SELECT C.NumEnrolled  
      INTO      :num_enrolled  
      FROM      Course C  
      WHERE     C.CrsCode = :crs_code;
```

# Parameters: Static vs. Dynamic SQL

- Dynamic SQL - SQL statement constructed at run time when symbol table is no longer present
- Case 1: Parameters are known at compile time

```
strcpy (tmp, "SELECT C.NumEnrolled FROM Course C  
          WHERE C.CrsCode = ?") ;  
EXEC SQL PREPARE st FROM :tmp;
```

- Parameters are named in EXECUTE statement: *in* parameters in USING; *out* parameters in INTO clauses

```
EXEC SQL EXECUTE st INTO :num_enrolled USING :crs_code;
```

- EXECUTE statement is compiled using symbol table
  - *fetch()* and *store()* routines generated

# Parameters - Dynamic SQL

(parameters known at compile time)

- Fetch and store routines are executed at client when EXECUTE is executed to communicate argument values with DBMS
- EXECUTE can be invoked multiple times with different values of *in* parameters
  - Each invocation uses same query execution plan
- Values substituted for placeholders by DBMS (in order) at invocation time and statement is executed

# Parameters in Dynamic SQL

(parameters supplied at runtime)

- Case 2: Parameters not known at compile time
- *Example*: Statement input from terminal
  - Application cannot parse statement and might not know schema, so it does not have any parameter information
- EXECUTE statement cannot name parameters in INTO and USING clauses

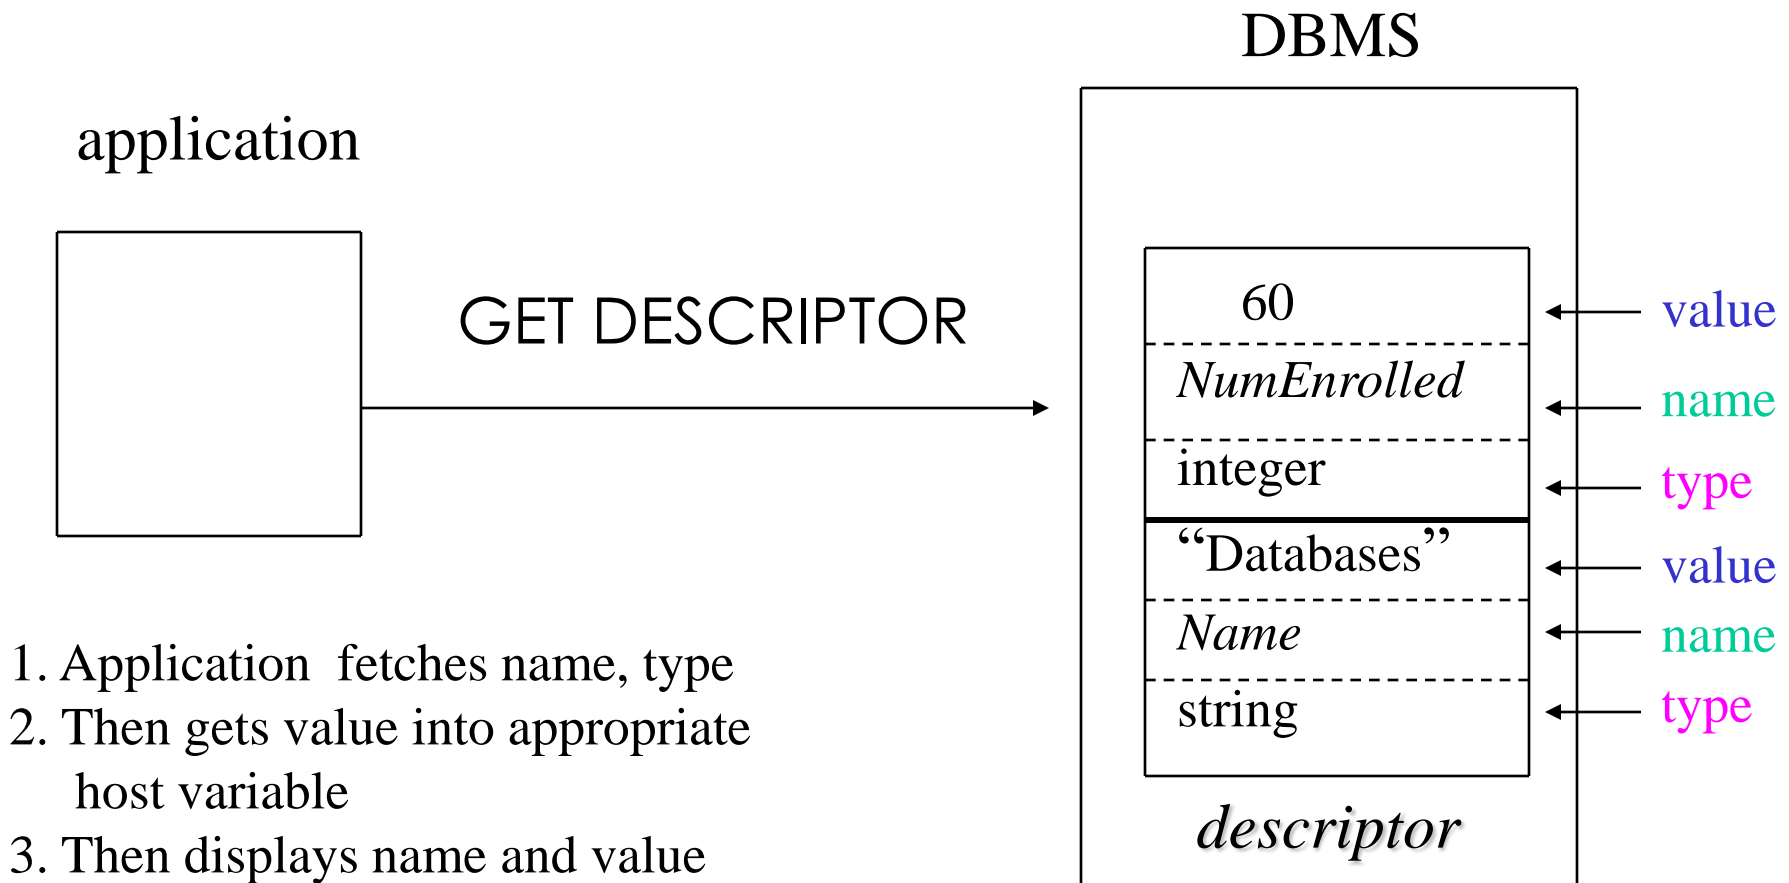
# Parameters in Dynamic SQL (cont'd)

## (parameters supplied at runtime)

- DBMS determines number and type of parameters after preparing the statement
- Information stored by DBMS in a *descriptor* – a data structure inside the DBMS, which records the *name*, *type*, and *value* of each parameter
- Dynamic SQL provides directive GET DESCRIPTOR to get information about parameters (e.g., number, name, type) from DBMS and to fetch value of *out* parameters
- Dynamic SQL provides directive SET DESCRIPTOR to supply value to *in* parameters

# Descriptors

```
temp = "SELECT C.NumEnrolled, C.Name FROM Course C \
WHERE C.CrsCode = 'CS305' "
```



# Dynamic SQL Calls when Descriptors are Used

... .. *construct SQL statement in temp* ... ..

```
EXEC SQL PREPARE st FROM :temp;           // prepare statement
```

```
EXEC SQL ALLOCATE DESCRIPTOR 'desc';      // create descriptor
```

```
EXEC SQL DESCRIBE OUTPUT st USING  
      SQL DESCRIPTOR 'desc';              // populate desc with info  
                                           // about out parameters
```

```
EXEC SQL EXECUTE st INTO                  // execute statement and  
      SQL DESCRIPTOR AREA 'desc';         // store out values in desc
```

```
EXEC SQL GET DESCRIPTOR 'desc' ...;       // get out values
```

... .. *similar strategy is used for in parameters* ... ..

# Example: Nothing Known at Compile Time

```
sprintf(my_sql_stmt,  
        "SELECT * FROM %s WHERE COUNT(*) = 1",  
        table); // table – host var; even the table is known only at run time!
```

```
EXEC SQL PREPARE st FROM :my_sql_stmt;  
EXEC SQL ALLOCATE DESCRIPTOR 'st_output';
```

```
EXEC SQL DESCRIBE OUTPUT st USING SQL DESCRIPTOR 'st_output'
```

- The SQL statement to execute is known only at run time
- At this point DBMS knows what the exact statement is (including the table name, the number of *out* parameters, their types)
- The above statement asks to create descriptors in `st_output` for all the (now known) *out* parameters

```
EXEC SQL EXECUTE st INTO SQL DESCRIPTOR 'st_output';
```



# Example: Getting Meta-Information from a Descriptor

- More details in the book:

```
// Host var colcount gets the number of out parameters in the SQL statement
```

```
// described by st_output
```

```
EXEC SQL GET DESCRIPTOR 'st_output' :colcount = COUNT;
```

```
// Set host vars coltype, collength, colname with the type, length, and name
```

```
// of the 3-d out parameter in the SQL statement described by st_output
```

```
EXEC SQL GET DESCRIPTOR 'st_output' VALUE 3;
```

```
:coltype = TYPE, // predefined integer constants, such as SQL_CHAR, SQL_FLOAT,...
```


```
:collength = LENGTH,
```

```
:colname = NAME;
```

# JDBC

- Call-level interface (CLI) for executing SQL from a Java program
- SQL statement is constructed at run time as the value of a Java variable (as in dynamic SQL)
- JDBC passes SQL statements to the underlying DBMS. Can be interfaced to any DBMS that has a JDBC driver

# JDBC Run-Time Architecture



# Executing a Query

```
import java.sql.*;    -- import all classes in package java.sql

Class.forName (driver name);    // static method of class Class
                                // loads specified driver

Connection con = DriverManager.getConnection(Url, Id, Passwd);
    // Static method of class DriverManager; attempts to
    // connect to DBMS; if successful, creates a connection
    // object, con, for managing the connection

Statement stat = con.createStatement ();
    // Creates a statement object stat;
    // statements have executeQuery() method
```

# Executing a Query

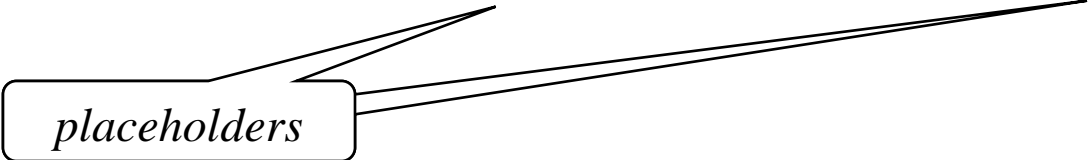
```
String query = "SELECT  T.StudId FROM Transcript T" +  
               "WHERE  T.CrsCode = 'cse305' " +  
               "AND  T.Semester = 'S2000' ";
```

```
ResultSet res = stat.executeQuery(query);
```

- *Creates a result set object, res.*
- *Prepares and executes the query.*
- *Stores the result set produced by execution in res (analogous to opening a cursor).*
- *The query string can be constructed at run time (as above).*
- *The input parameters are plugged into the query when the string is formed (as above)*

# Preparing and Executing a Query

```
String query = "SELECT T.StudId FROM Transcript T" +  
               "WHERE T.CrsCode = ? AND T.Semester = ?";
```



*placeholders*

```
PreparedStatement ps = con.prepareStatement ( query );
```

- *Prepares the statement*
- *Creates a prepared statement object, ps, containing the prepared statement*
- *Placeholders (?) mark positions of in parameters; special API is provided to plug the actual values in positions indicated by the ?'s*

# Preparing and Executing a Query

```
String crs_code, semester;
```

```
.....
```

```
ps.setString(1, crs_code);    // set value of first in parameter
```

```
ps.setString(2, semester);    // set value of second in parameter
```

```
ResultSet res = ps.executeQuery ( );
```

```
    // Creates a result set object, res.
```

```
    // Executes the query
```

```
    // Stores the result set produced by execution in res.
```

```
while (res.next ( ) ) {
```

```
    // advance the cursor
```

```
    j = res.getInt (“StudId”);
```

```
    // fetch output int-value
```

```
    ...process output value...
```

```
}
```

# Result Sets and Cursors

- Three types of result sets in JDBC:
  - *Forward-only*: not scrollable
  - *Scroll-insensitive*: scrollable, changes made to underlying tables after the creation of the result set are not visible through that result set
  - *Scroll-sensitive*: scrollable, changes made to the tuples in a result set after the creation of that set are visible through the set



# Result Set

```
Statement stat = con.createStatement (
    ResultSet.TYPE_SCROLL_SENSITIVE,
    ResultSet.CONCUR_UPDATABLE );
```

- Any result set type can be declared *read-only* or *updatable* – **CONCUR\_UPDATABLE** (assuming SQL query satisfies the conditions for updatable views)
- Current row of an updatable result set can be updated or deleted, and a new row can be inserted, causing changes in base table

```
res.updateString ("Name", "John" );    // update attribute "Name" of
                                         // current row in row buffer.
```

```
res.updateRow ( );    // install current row buffer in res and underlying table
```

# Handling Exceptions

```
try {  
    ...Java/JDBC code...  
} catch ( SQLException ex ) {  
    ...exception handling code...  
}
```

- try/catch is the basic structure within which an SQL statement should be embedded
- If an exception is thrown, an exception object, *ex*, is created and the catch clause is executed
- The exception object has methods to print an error message, return SQLSTATE, etc.

# Transactions in JDBC

- Default for a connection is
  - Transaction boundaries
    - *Autocommit mode*: each SQL statement is a transaction.
    - To group several statements into a transaction use `con.setAutocommit(false)`
  - Isolation
    - default isolation level of the underlying DBMS
    - To change isolation level use `con.setTransactionIsolationLevel(TRANSACTION_SERIALIZABLE)`
- With autocommit off:
  - transaction is committed using `con.commit()`.
  - next transaction is automatically initiated (chaining)
- Transactions on each connection committed separately

# SQLJ

- A statement-level interface to Java
  - A dialect of embedded SQL designed specifically for Java
  - Translated by precompiler into Java
  - SQL constructs translated into calls to an SQLJ runtime package, which accesses database through calls to a JDBC driver

# SQLJ

- Has some of efficiencies of embedded SQL
  - Compile-time syntax and type checking
  - Use of host language variables
  - More elegant than embedded SQL
- Has some of the advantages of JDBC
  - Can access multiple DBMSs using drivers
  - SQLJ statements and JDBC calls can be included in the same program

# SQLJ Example

```
#SQL {  
    SELECT C.Enrollment  
    INTO :numEnrolled  
    FROM Class C  
    WHERE C.CrsCode = :crsCode  
           AND C.Semester = :semester  
};
```

# Example of SQLJ Iterator

```
#SQL iterator GetEnrolledIter (int studentId, String studGrade);  
GetEnrolledIter iter1;
```

```
#SQL iter1 = {  
    SELECT T.StudentId as “studentId”,  
           T.Grade as “studGrade”  
    FROM Transcript T  
    WHERE T.CrsCode = :crsCode  
           AND T.Semester = :semester  
};
```

*Method names by  
which to access the  
attrs StudentId,  
Grade*

# Iterator Example (Cont'd)

```
int id;  
String grade;  
while (iter1.next() {  
    id = iter1.studentId();  
    grade = iter1.studGrade();  
    ... process the values in id and grade ...  
};  
  
iter1.close();
```



# ODBC

- Call level interface that is database independent
- Related to SQL/CLI, part of SQL:1999
- Software architecture similar to JDBC with driver manager and drivers
- Not object oriented
- Low-level: application must specifically allocate and deallocate storage

# Sequence of Procedure Calls Needed for ODBC

```
SQLAllocEnv(&henv);           // get environment handle
SQLAllocConnect(henv, &hdbc); // get connection handle
SQLConnect(hdbc, db_name, userId, password); // connect
SQLAllocStmt(hdbc, &hstmt);   // get statement handle
SQLPrepare(hstmt, SQL statement); // prepare SQL statement
SQLExecute(hstmt);
SQLFreeStmt(hstmt);           // free up statement space
SQLDisconnect(hdbc);
SQLFreeEnv(henv);             // free up environment space
```

# ODBC Features

- Cursors
  - *Statement handle* (for example hstmt) is used as name of cursor
- Status Processing
  - Each ODBC procedure is actually a function that returns status

```
RETCODE retcode1;  
Retcode1 = SQLConnect ( ...)
```
- Transactions
  - Can be committed or aborted with

```
SQLTransact (henv, hdbc, SQL_COMMIT)
```